

# Actor Clustering with Docker Containers and Akka.NET

# Actor Clustering with Docker Containers and Akka.NET

# *“Why should we use the Actor programming model?”*

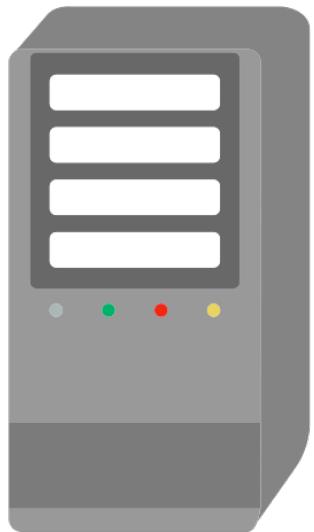
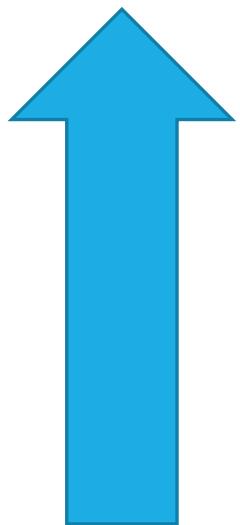
- Fault Tolerant
- Elastic
- Decentralized
- Scalable



# Its about maximizing resource use

---

**Scale up**



**vs**

**Scale out**

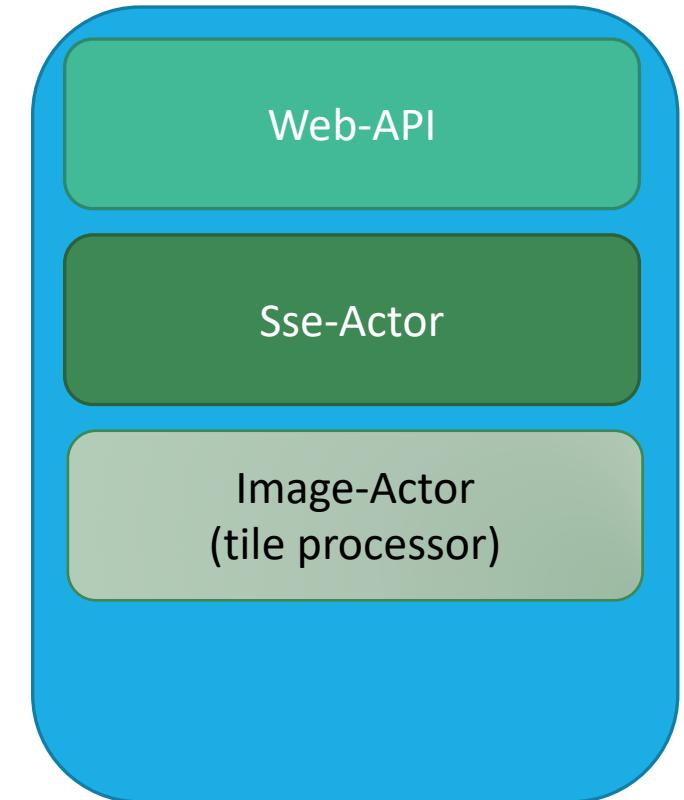
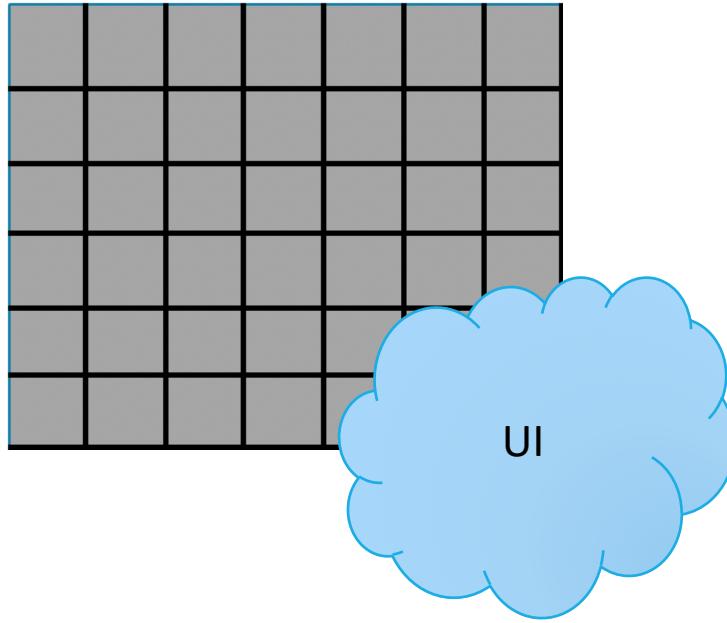


# What to expect

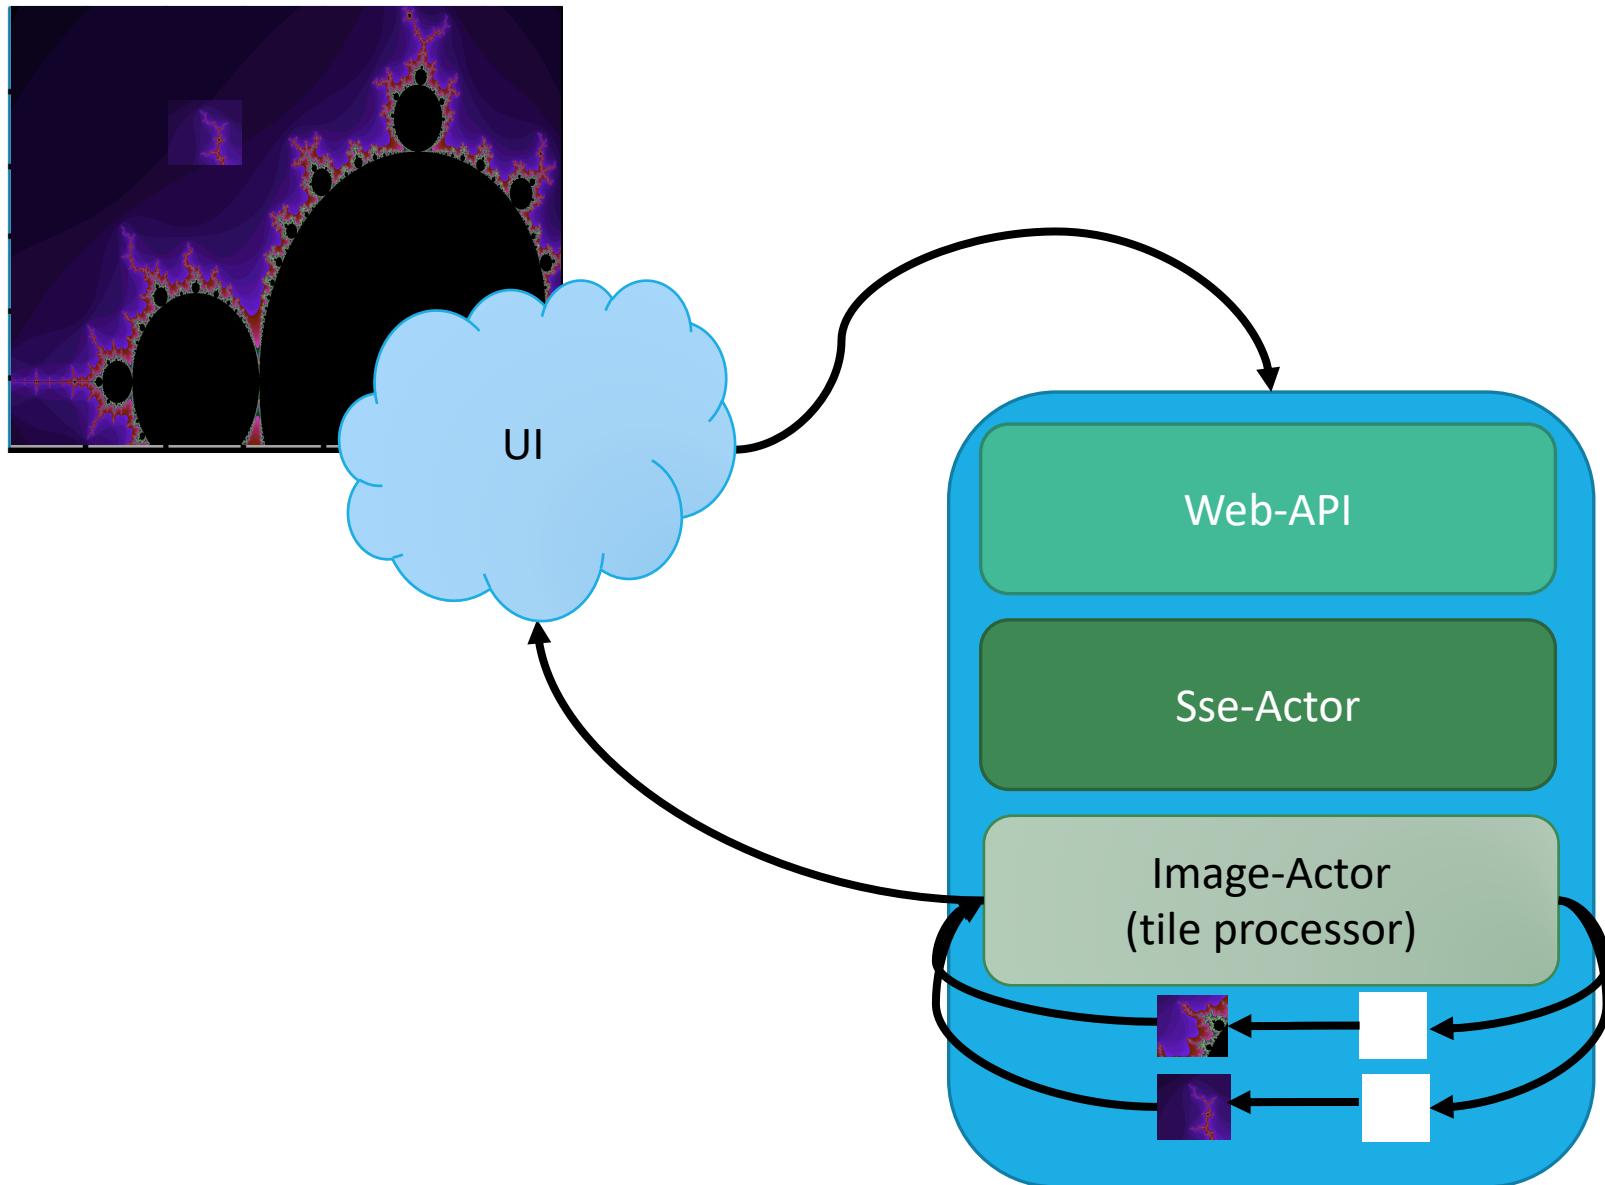


- **Akka** – core actor library
- **Akka.Persistence** – event-sourcing, durable actor state & recovery
- **Akka.Remote** – cross-node actor deployment / communication
- **Akka.Cluster** – elastic actor networks, distributed system
- **Docker** - platform-as-a-service products that use OS-level virtualization to deliver software in packages called containers.

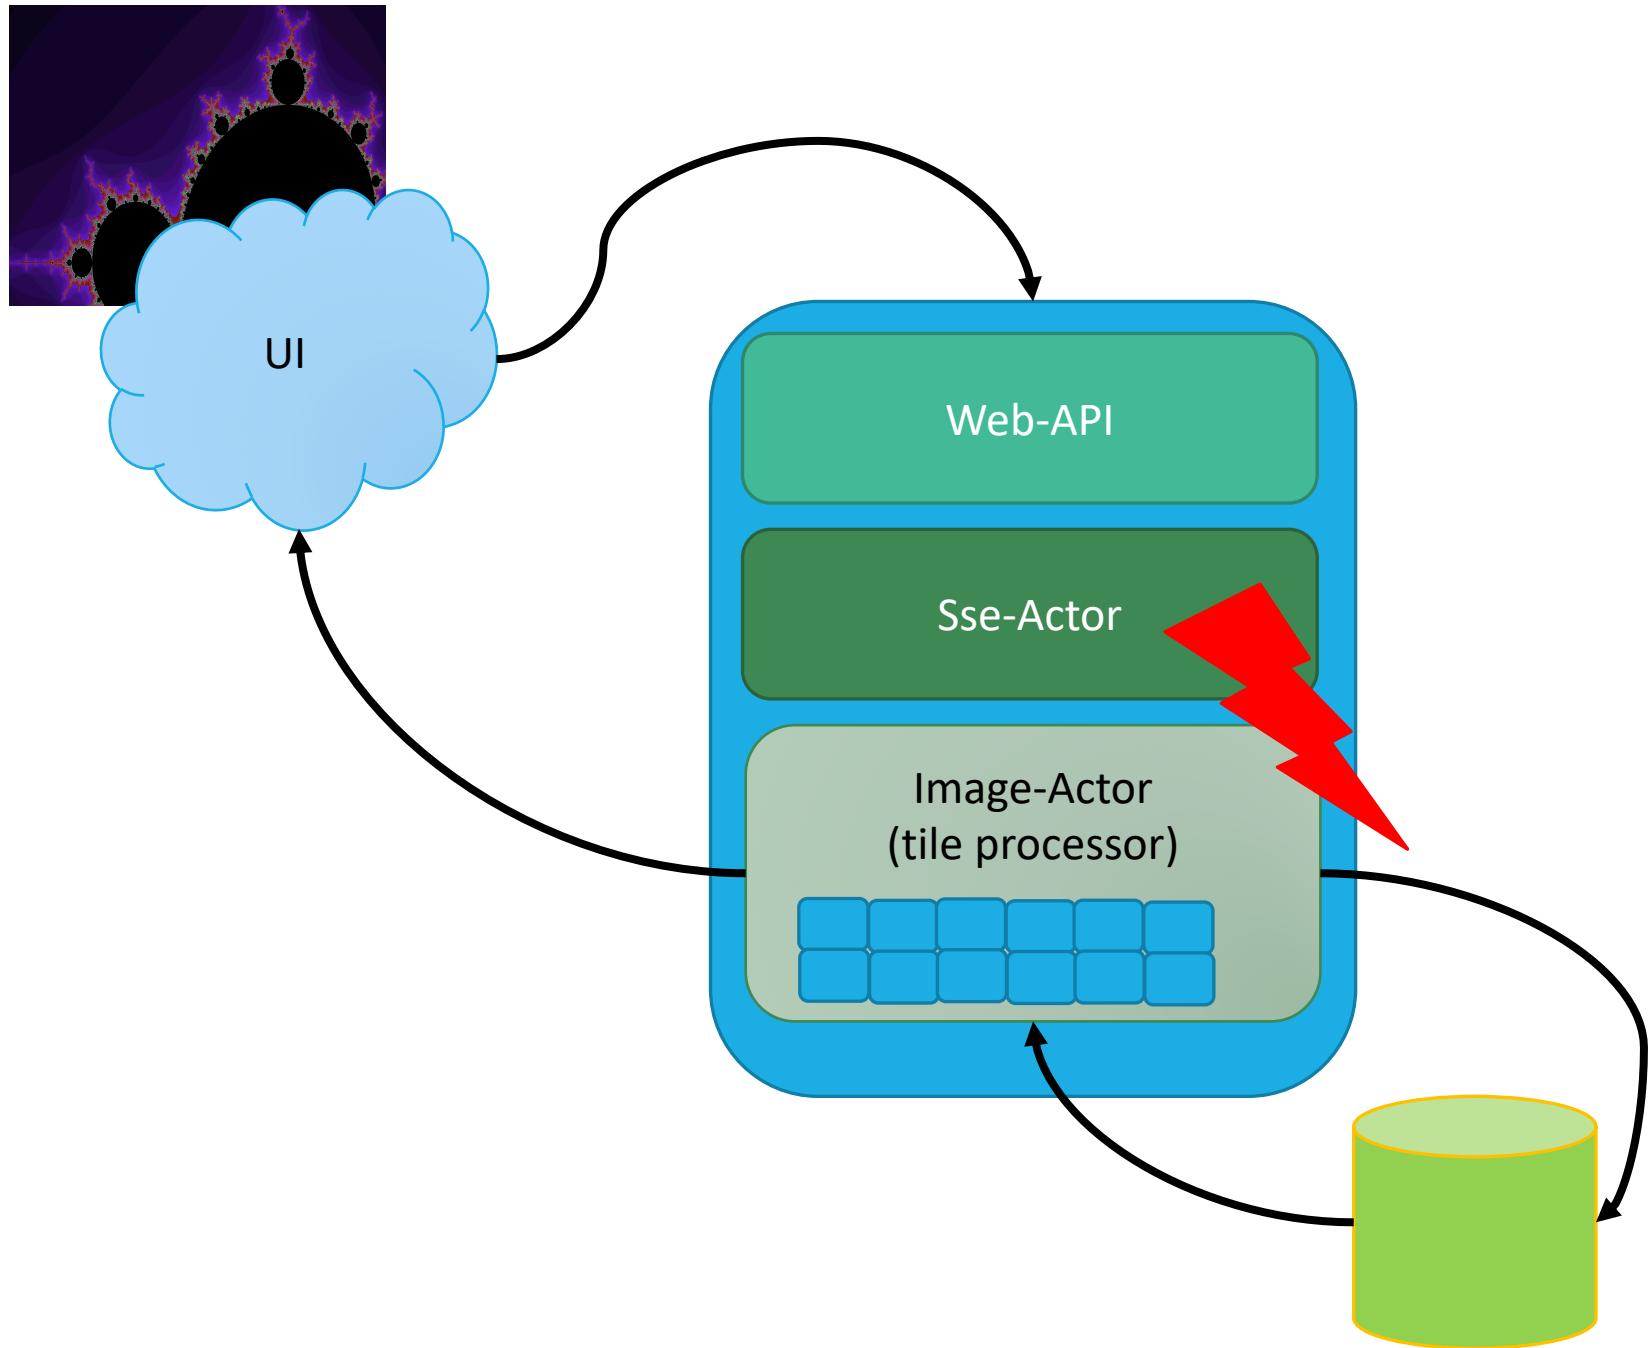
# Project Remote Distributed Image (fractal) Processing



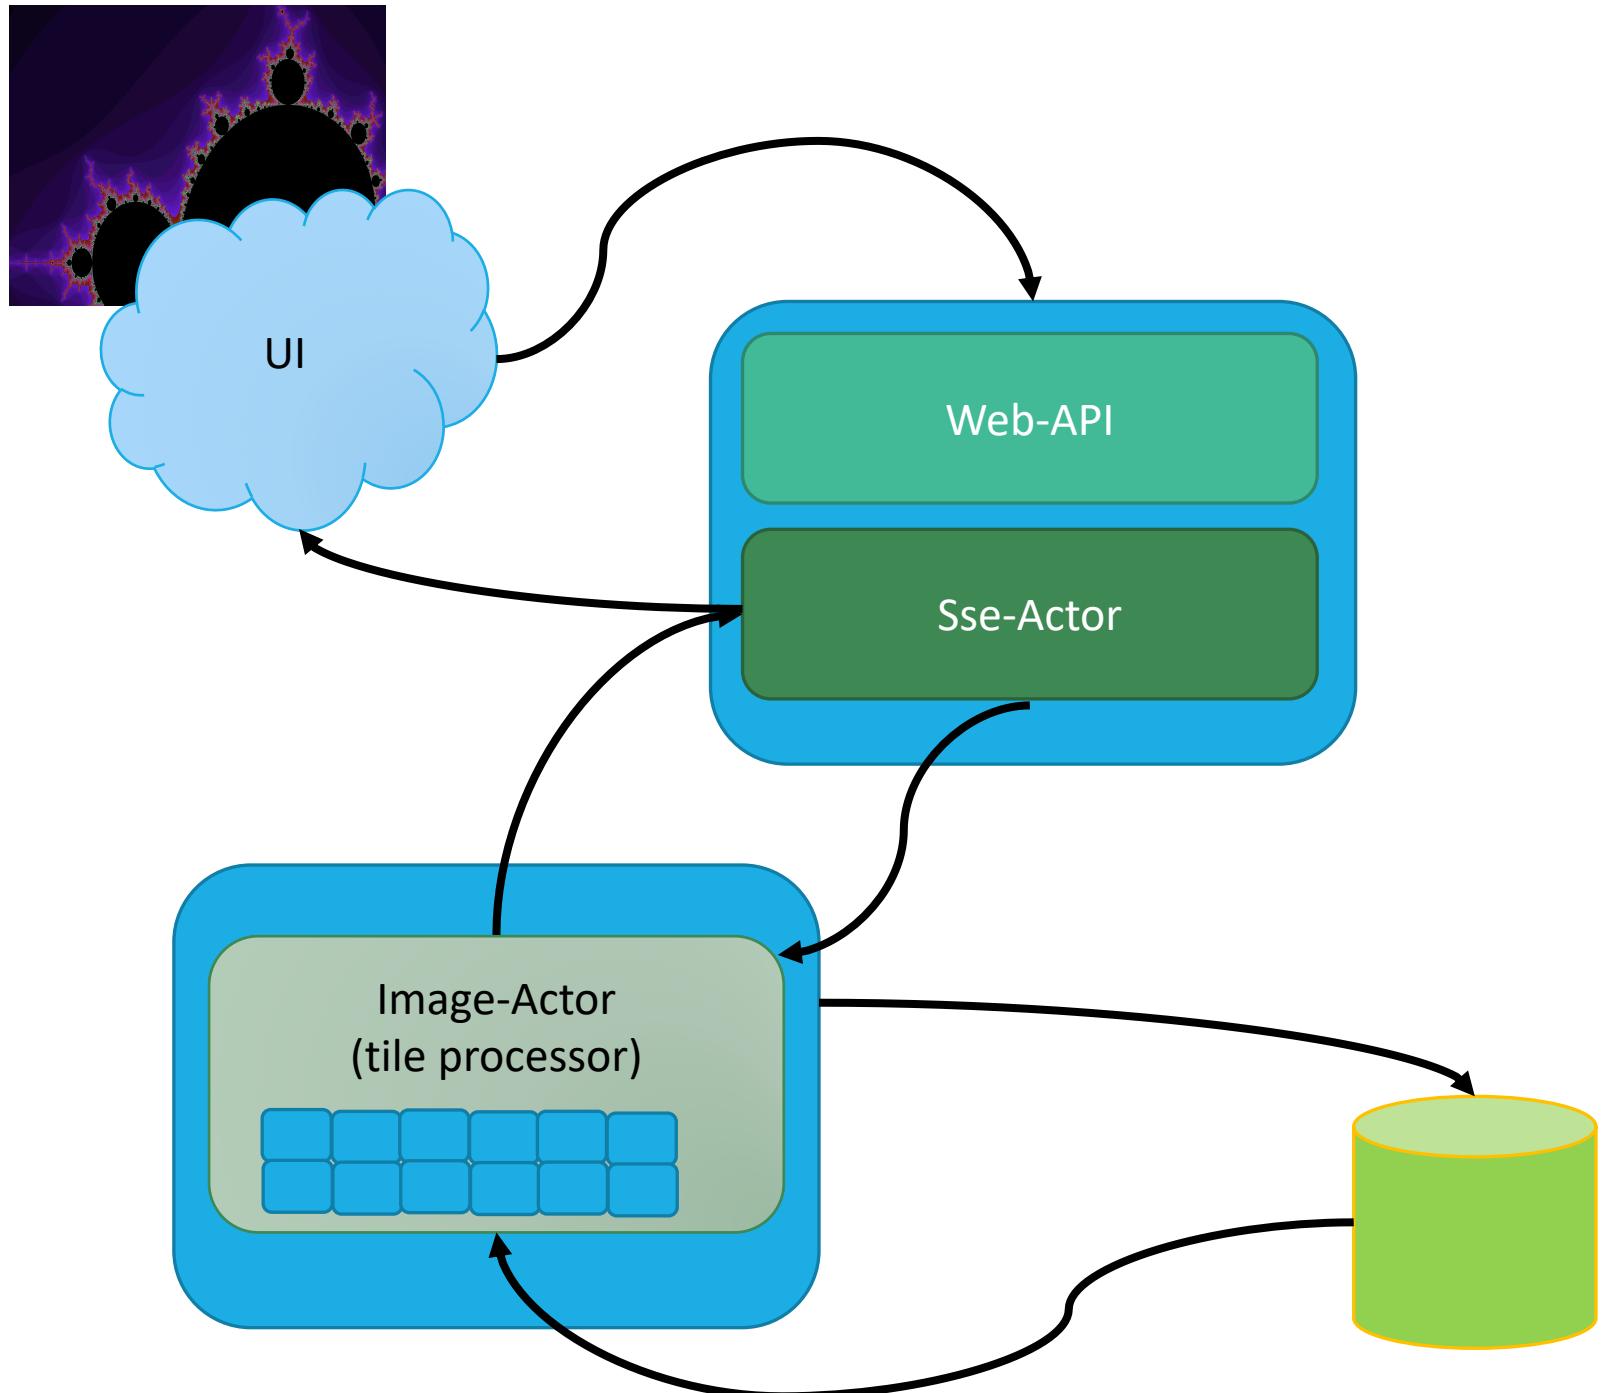
# Project Remote Distributed Image (fractal) Processing



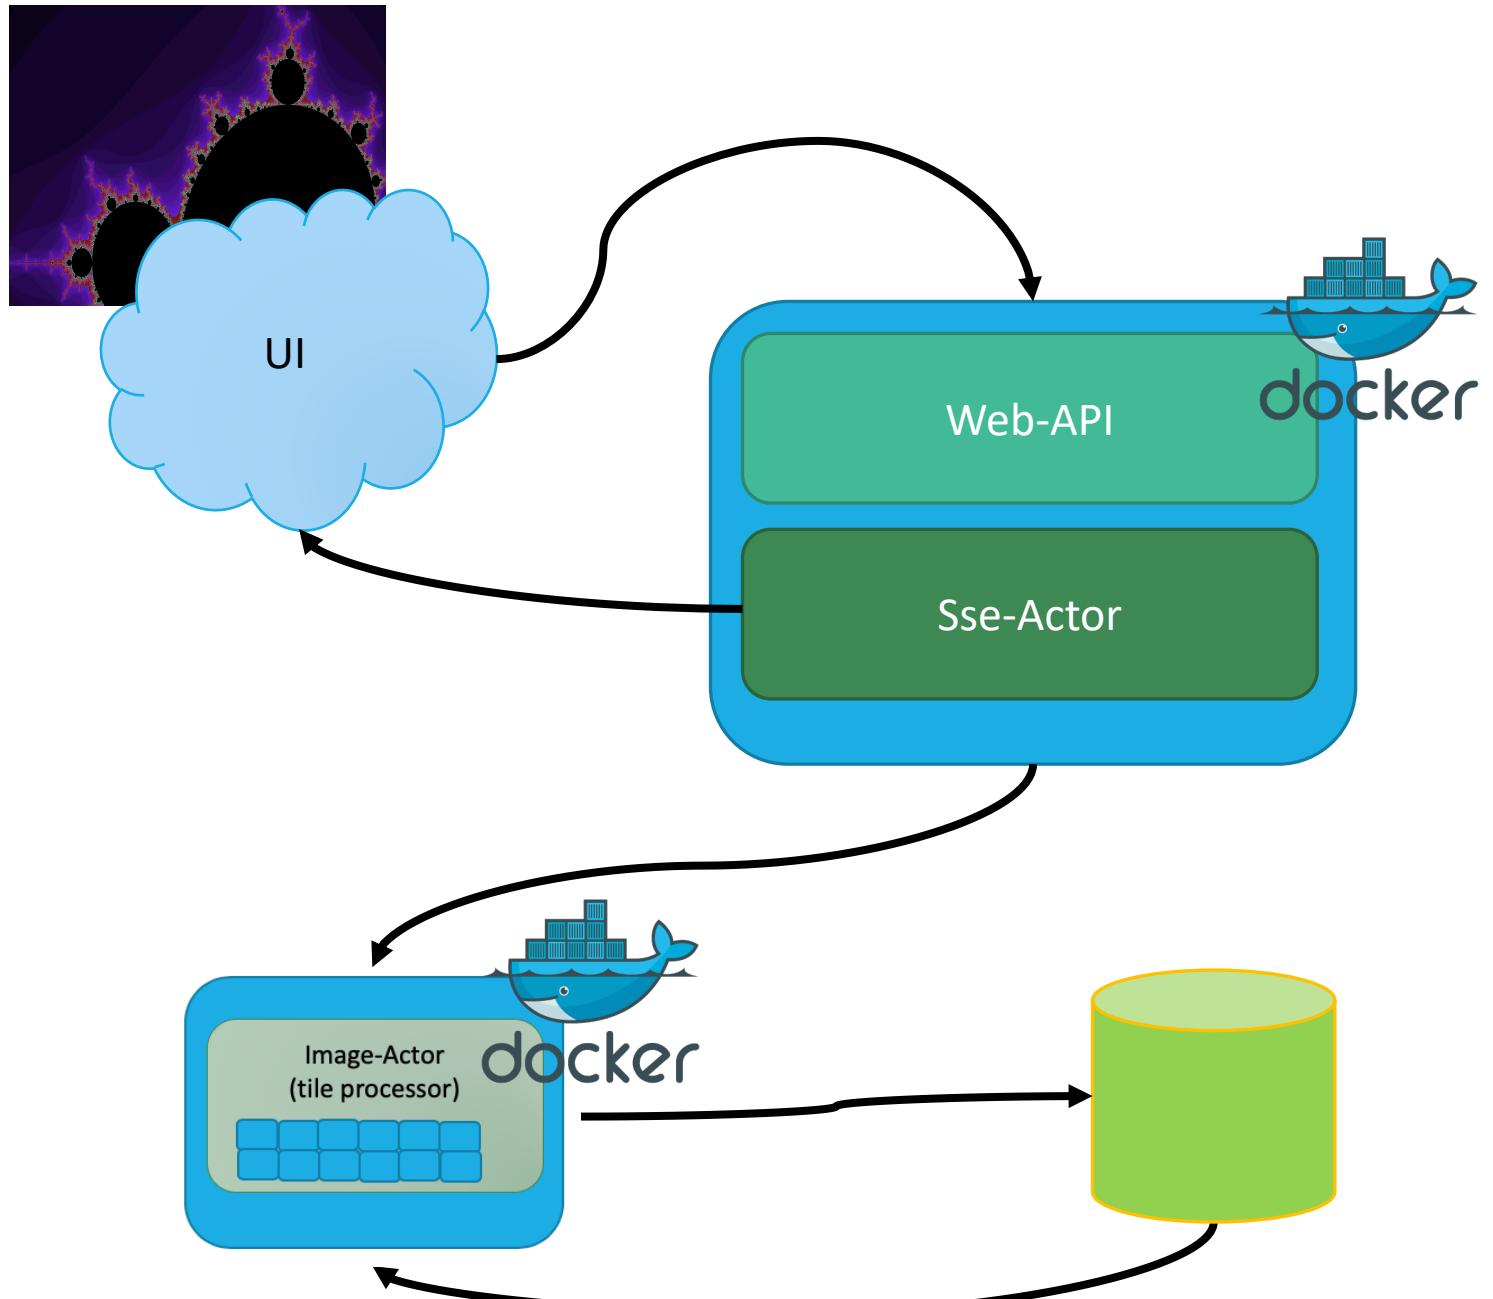
# Project Remote Distributed Image (fractal) Processing



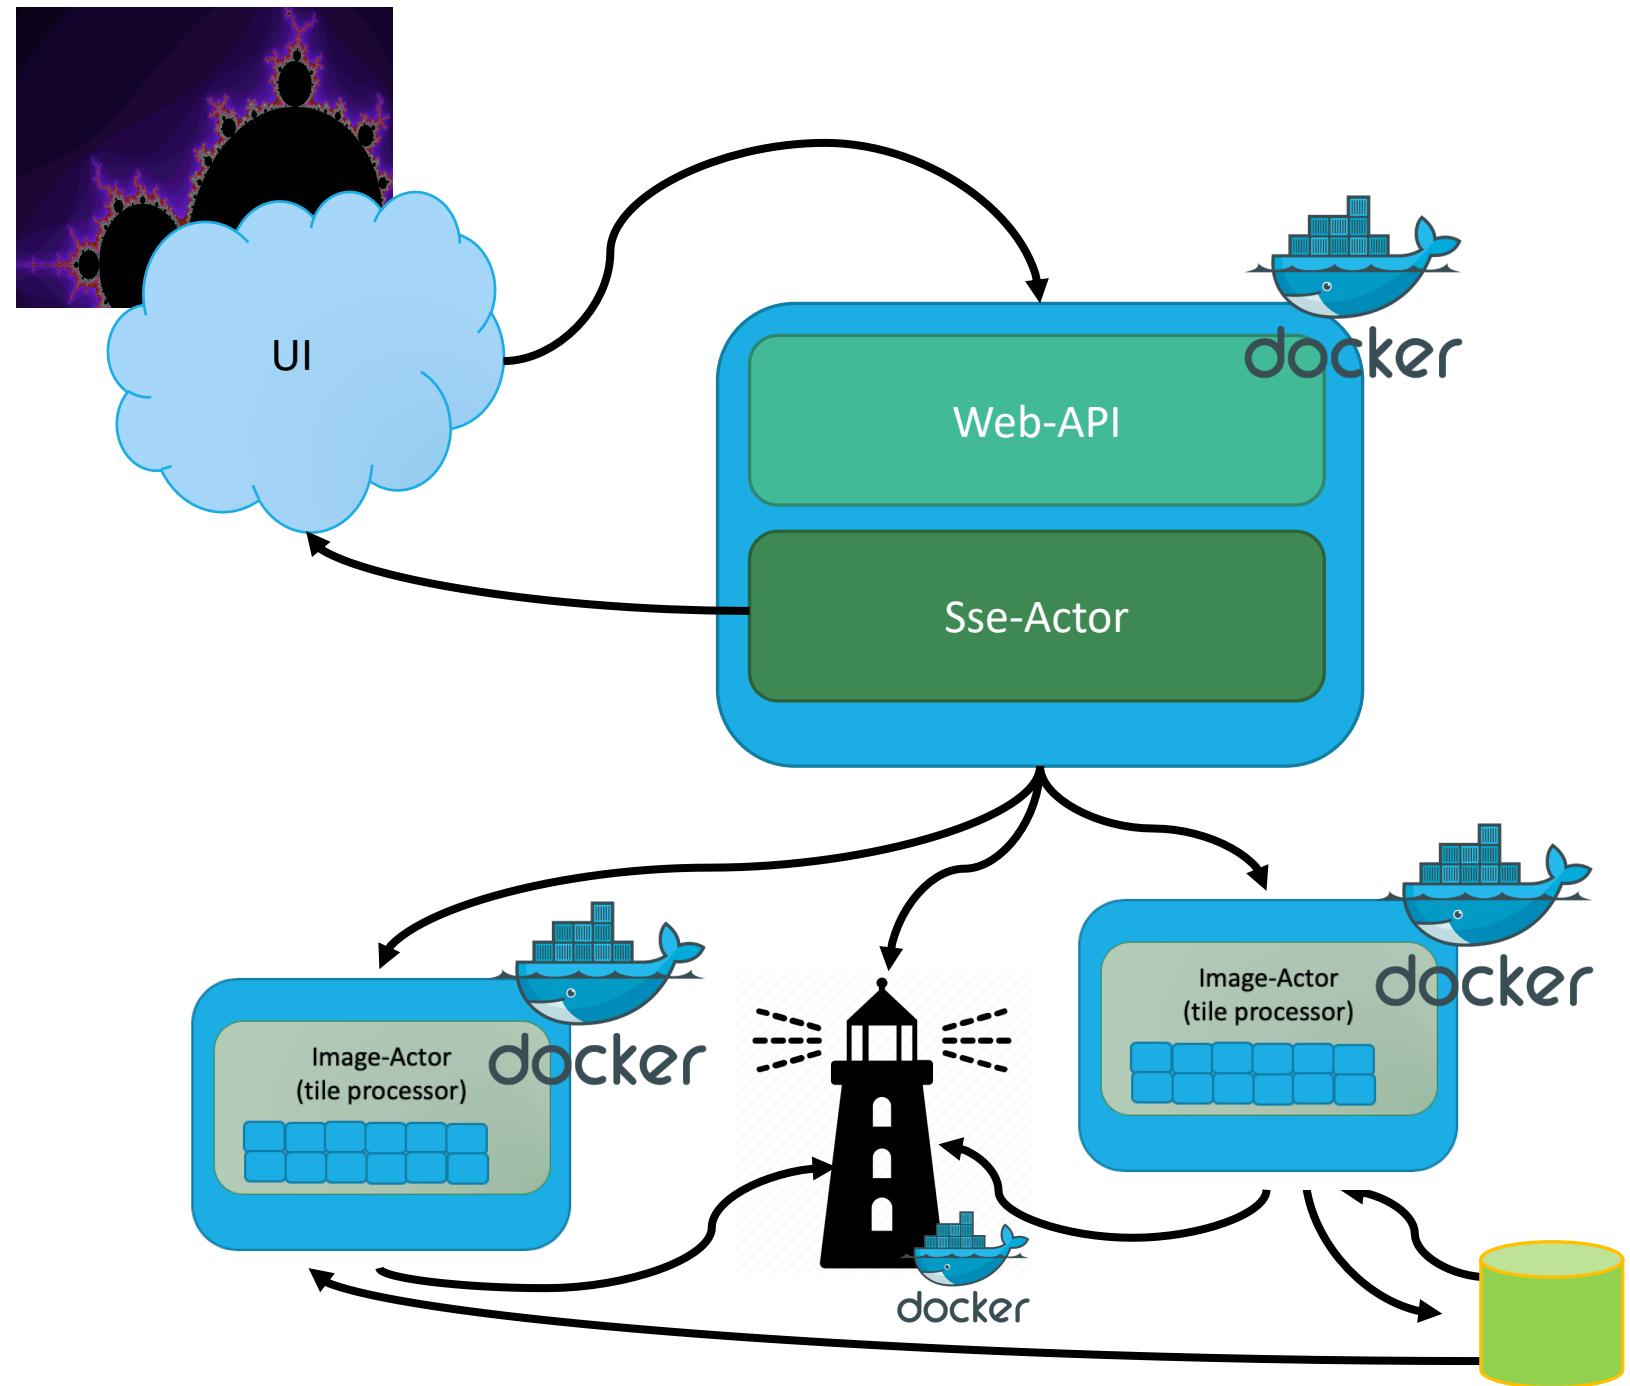
# Project Remote Distributed Image (fractal) Processing



# Project Remote Distributed Image (fractal) Processing

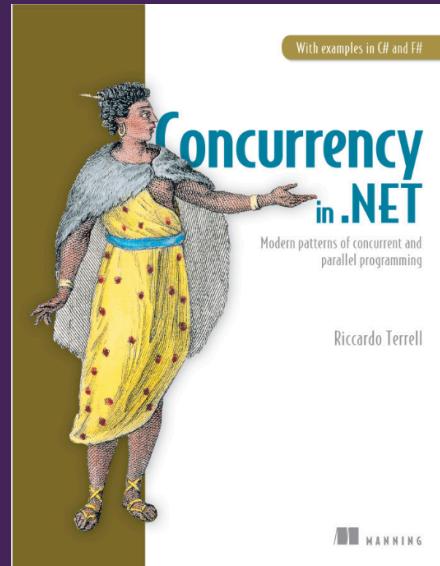


# Project Remote Distributed Image (fractal) Processing



# Introduction - Riccardo Terrell

- ④ Originally from Italy, currently - Living/working in Charlotte NC
- ④ +/- 20 years in professional programming
  - ④ C++/VB → Java → .Net C# → Scala → Haskell → C# & F# → ??
- ④ Author of the book “Concurrency in .NET” – Manning Pub.
- ④ *Polyglot programmer*
  - ④ *believes in the art of finding the right tool for the job*
- ④ *Organizer of the Pure Functional User Group*



# **MODERN PATTERNS OF PARALLEL AND DISTRIBUTED SYSTEMS IN .NET**

**2 DAY COURSE**

**THURSDAY, DEC. 5, 2019  
AT 9:00 AM-FRIDAY, DEC.  
6, 2019 AT 5:00 PM (EST)  
RALEIGH, NC**

**Next course in Raleigh  
December 5 and 6  
use code RaleighCC  
for 10% discount**

**Course link  
<https://bit.ly/2GeC3zD>**

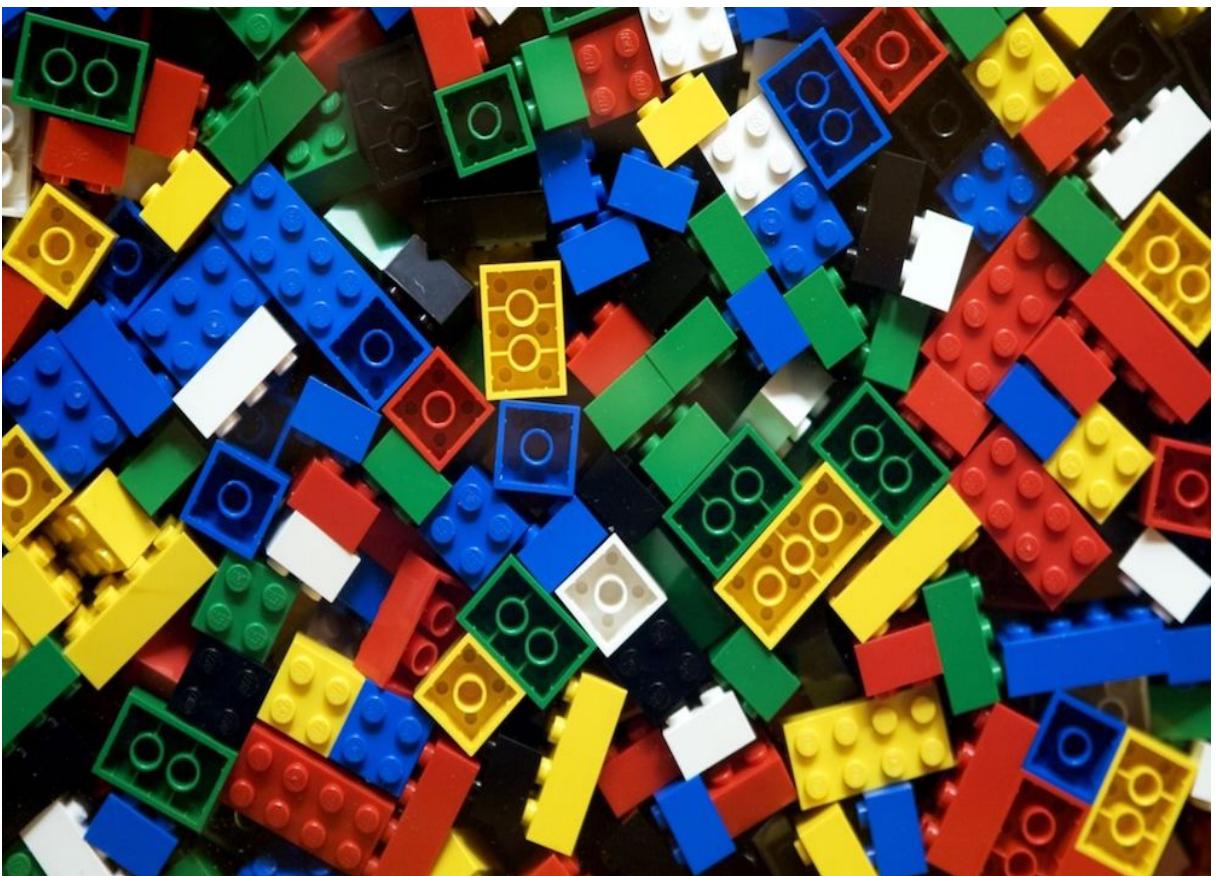
**For questions  
email  
[info@techstudiosoftware.com](mailto:info@techstudiosoftware.com)**

# Strategies to parallelize the code

---

The first step in designing any parallelized system is Decomposition.

Decomposition is nothing more than taking a problem space and breaking it into discrete parts. When we want to work in parallel, we need to have at least two separate things that we are trying to run. We do this by taking our problem and decomposing it into parts.



# The issue is Shared of Memory

---



Shared Memory Concurrency

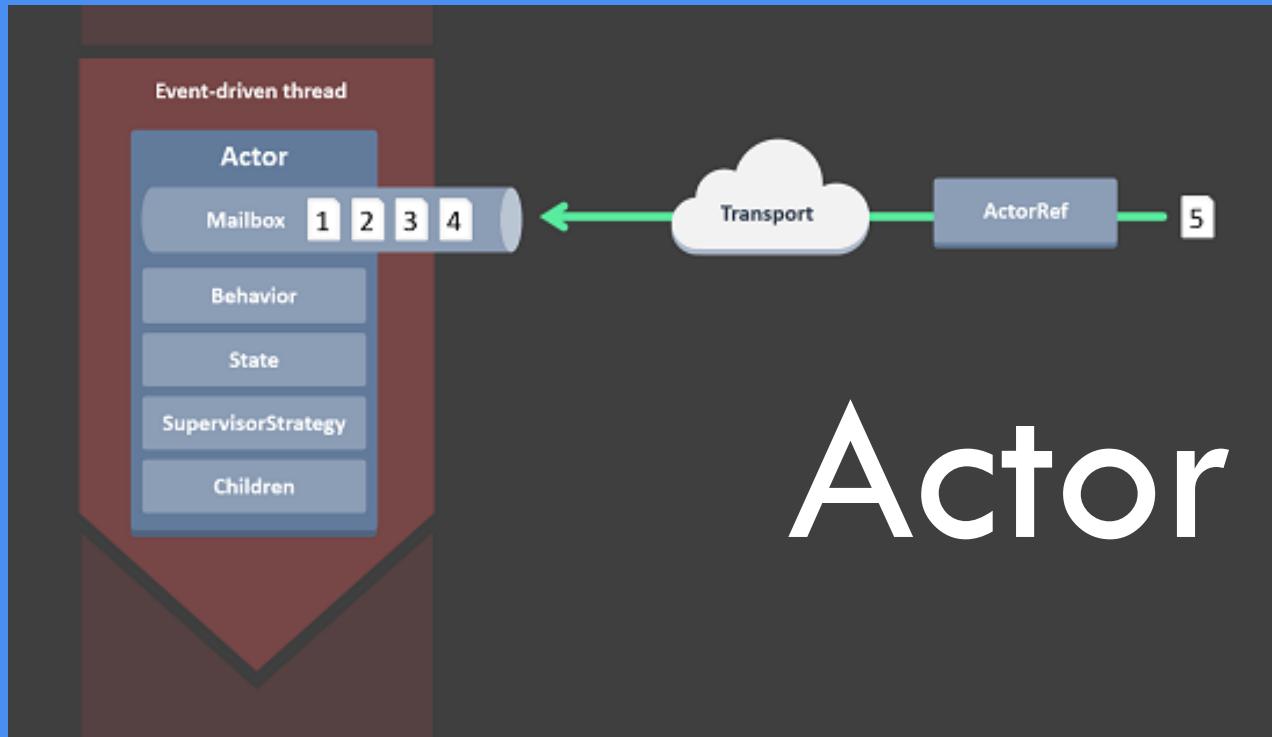
Data Race / Race Condition

Works in sequential single threaded environment

Not fun in a multi-threaded environment

Not fun trying to parallelize

Locking, blocking, call-back hell



# Actor Model

## Actor Model Three axioms:

1. Send messages to other Actors  
- *One Actor is not Actor* -
2. Create other Actor
3. Decide how to handle the next message

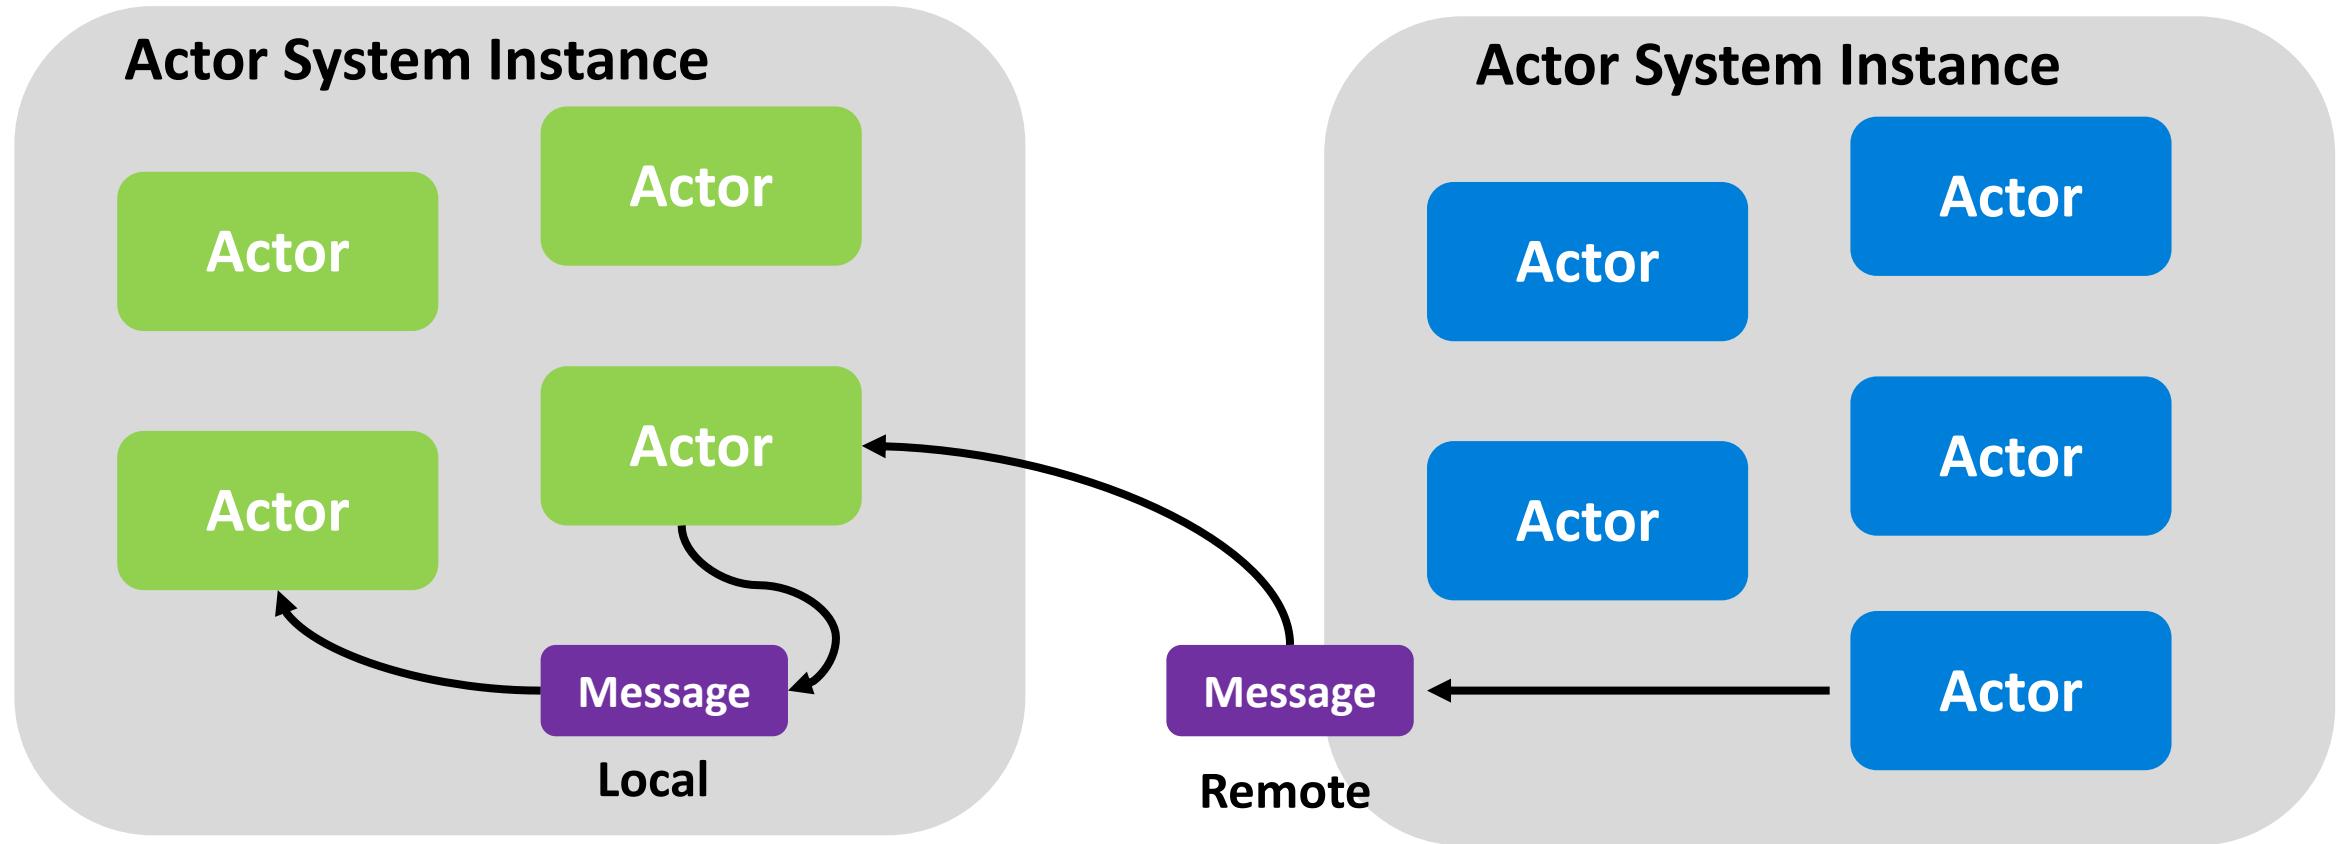
# What is an Actor?



- Processing
- Storage – State
- Communication only by messages
- Share Nothing
- Message are passed by value
- Lightweight object.
- Running on it's own thread.
- No shared state.
- Messages are kept in mailbox and processed in order.
- Massive scalable and lightening fast because of the small call stack.

# Actor Systems

---



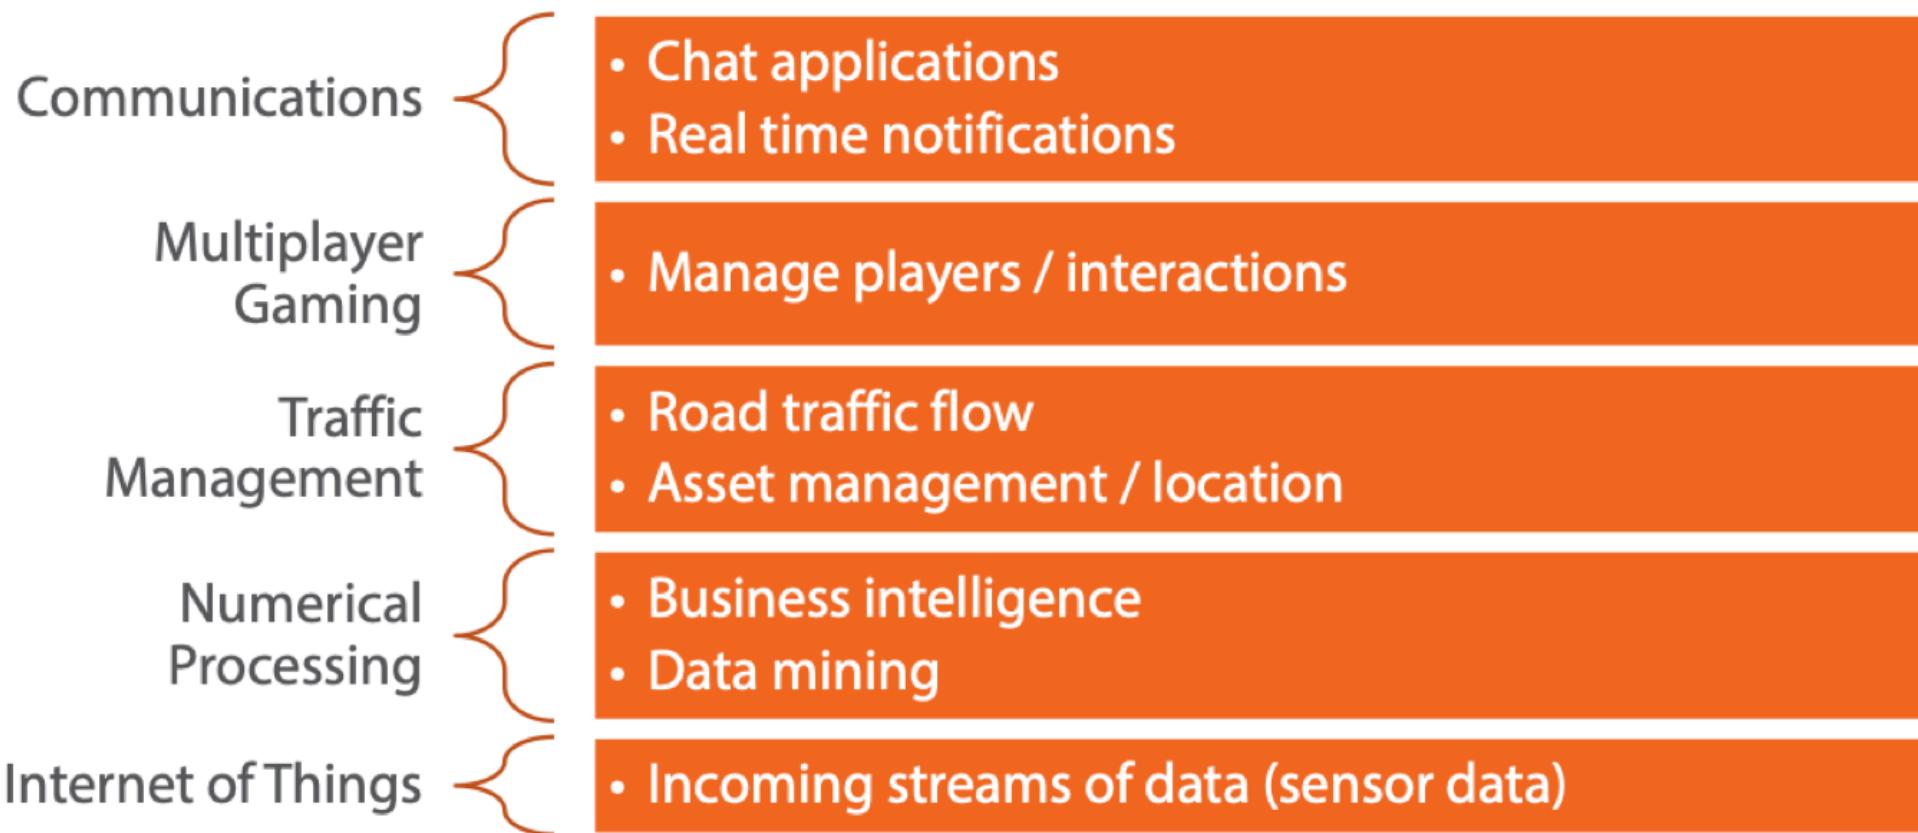
# Classes of Applications

---



# Classes of Applications

---



# Akka.NET



- Akka.Remote
- Akka.Cluster
- Akka.Persistence
  
- Supervision
- Routing

- \* *Concurrent*
- \* *Resilient*
  
- \* *Distributed*
- \* *Scalable*

# Key Features of Akka.NET

Simple Concurrency	<ul style="list-style-type: none"><li>• High level abstractions (actors, messages, FSMs)</li><li>• Asynchronous</li></ul>
Simple Distribution	<ul style="list-style-type: none"><li>• Location transparency</li><li>• Remote deployment using configuration</li></ul>
High Performance	<ul style="list-style-type: none"><li>• ~50 million messages per second per machine</li><li>• 1GB heap memory = ~2.5 million actors</li><li>• Load balancing / routing</li></ul>
Resilient	<ul style="list-style-type: none"><li>• Self-healing</li><li>• Supervisory hierarchies and fault isolation</li></ul>



# Actor – Akka.NET in C#

```
var system = ActorSystem.Create("fizz-buzz");

public class FizzBuzzActor : ReceiveActor {
    public FizzBuzzActor() {
        Receive<FizzBuzzMessage>(msg => {
            // code to handle the message
        });
    }
}

var actor = system.ActorOf(Props.Create<FizzBuzzActor>(), "fb-actor");
actor.Tell(new FizzBuzzMessage(5));
```



# Actor – Akka.NET in F#

```
let system = ActorSystem.Create("FSharp")

type EchoServer =
    inherit Actor

    override x.OnReceive (message:obj) =
        match message with
        | :? string as msg -> printfn "Hello %s" msg
        | _ -> printfn "What should I do with this thing?"

let echoServer = system.ActorOf(Props(typeof<EchoServer>))
echoServer.Tell 42
```

# Types of Message Sending

---

## Tell

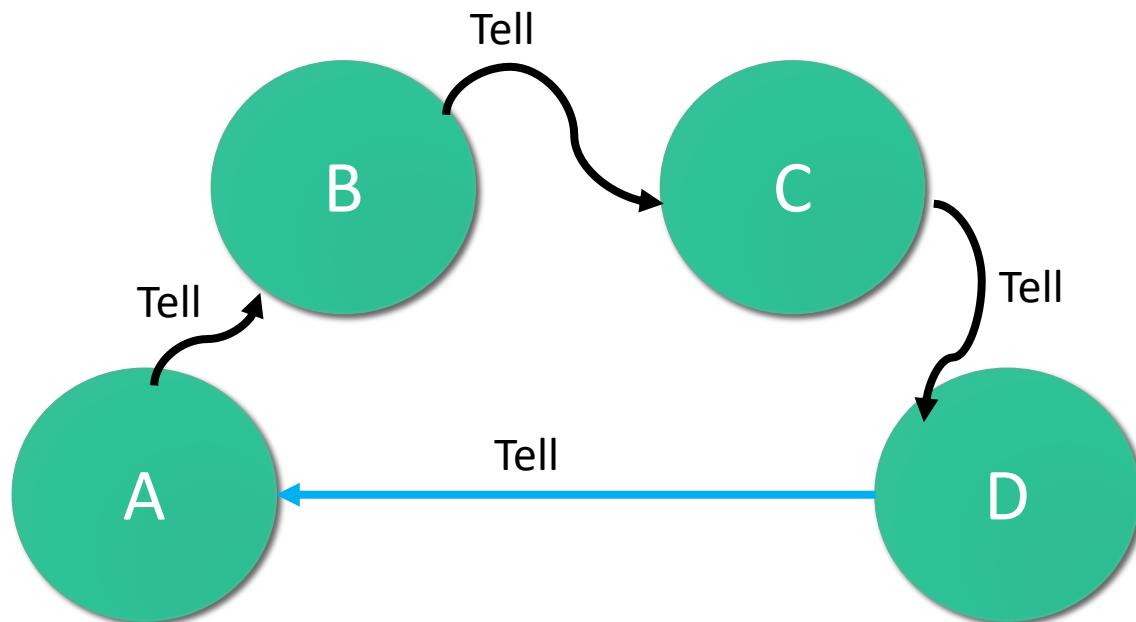
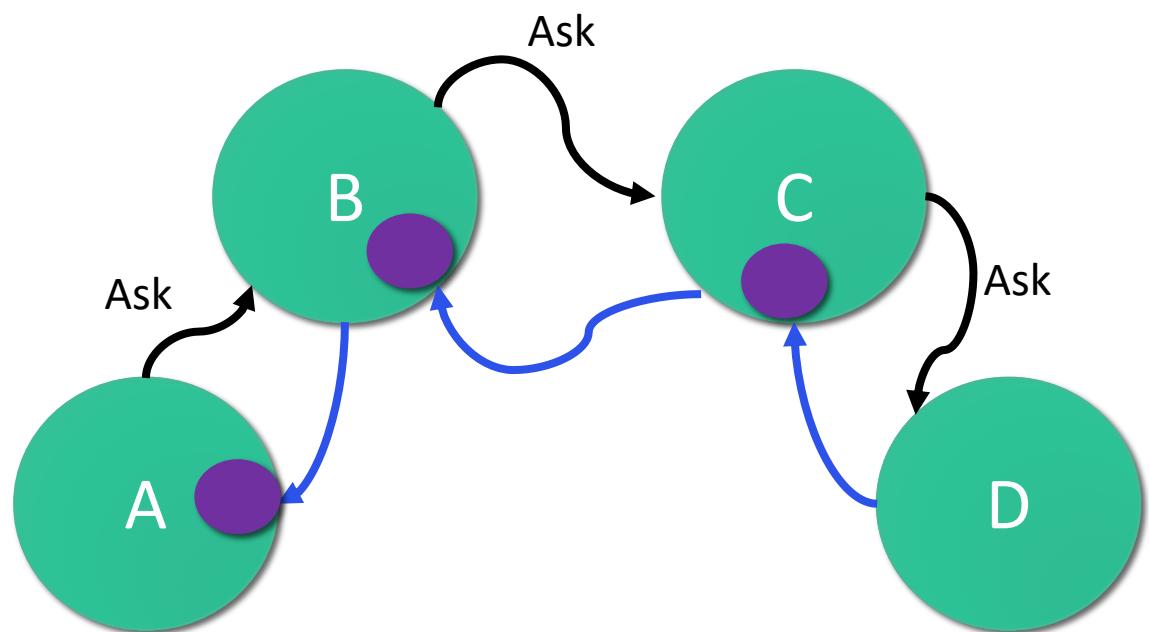
- Tell target actor something
- Don't expect a response
- "Fire and forget"
- Doesn't block waiting for a reply
- Better scaling and concurrency
- Use most of the time

## Ask

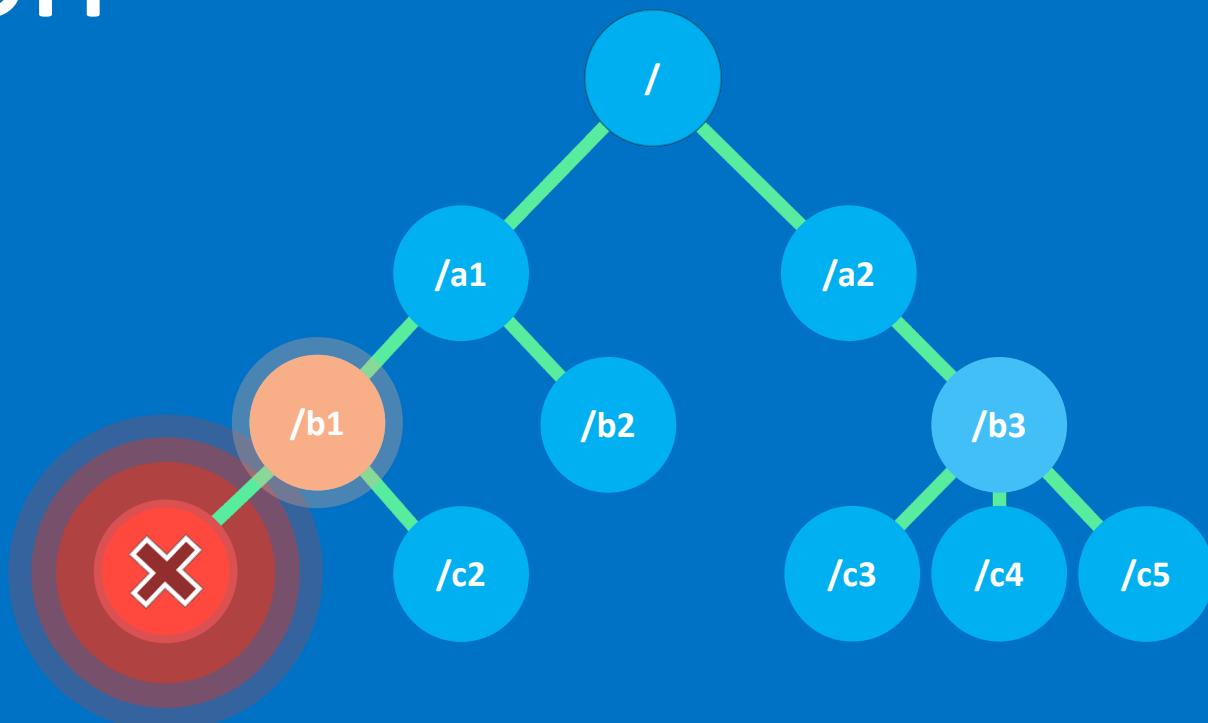
- Ask target actor for something
- Expect a response
- Target actor has to reply to sender
- Timeout (TaskCancelledException)
- Worse scaling and concurrency
- Only use when absolutely necessary

# Don't Ask, Tell

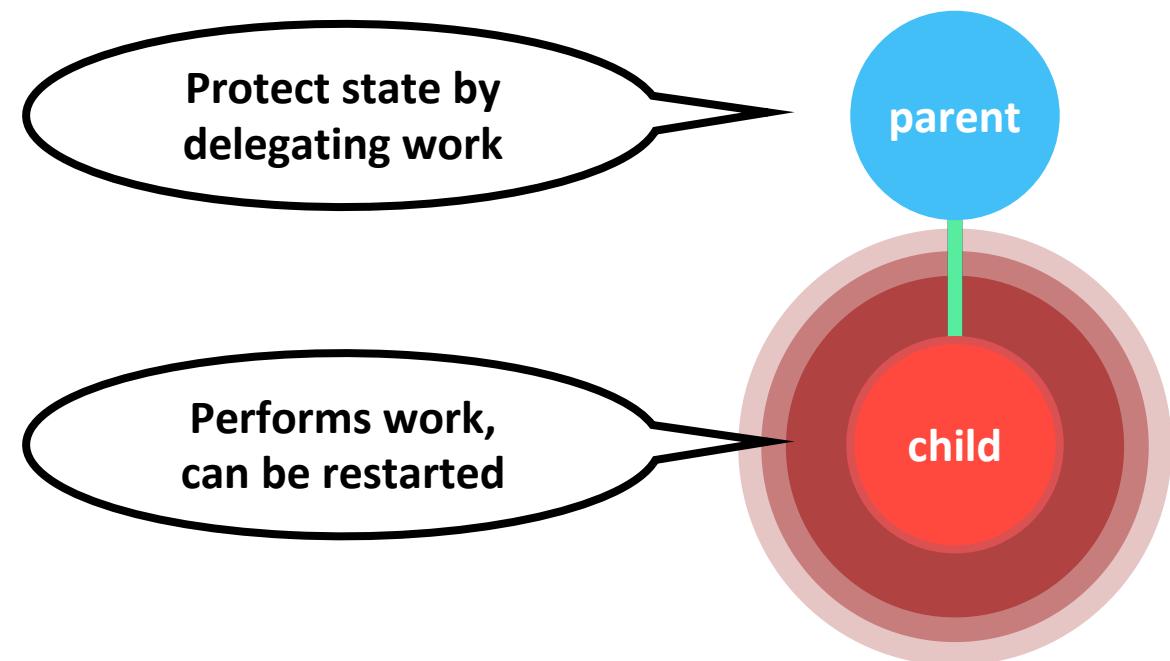
---

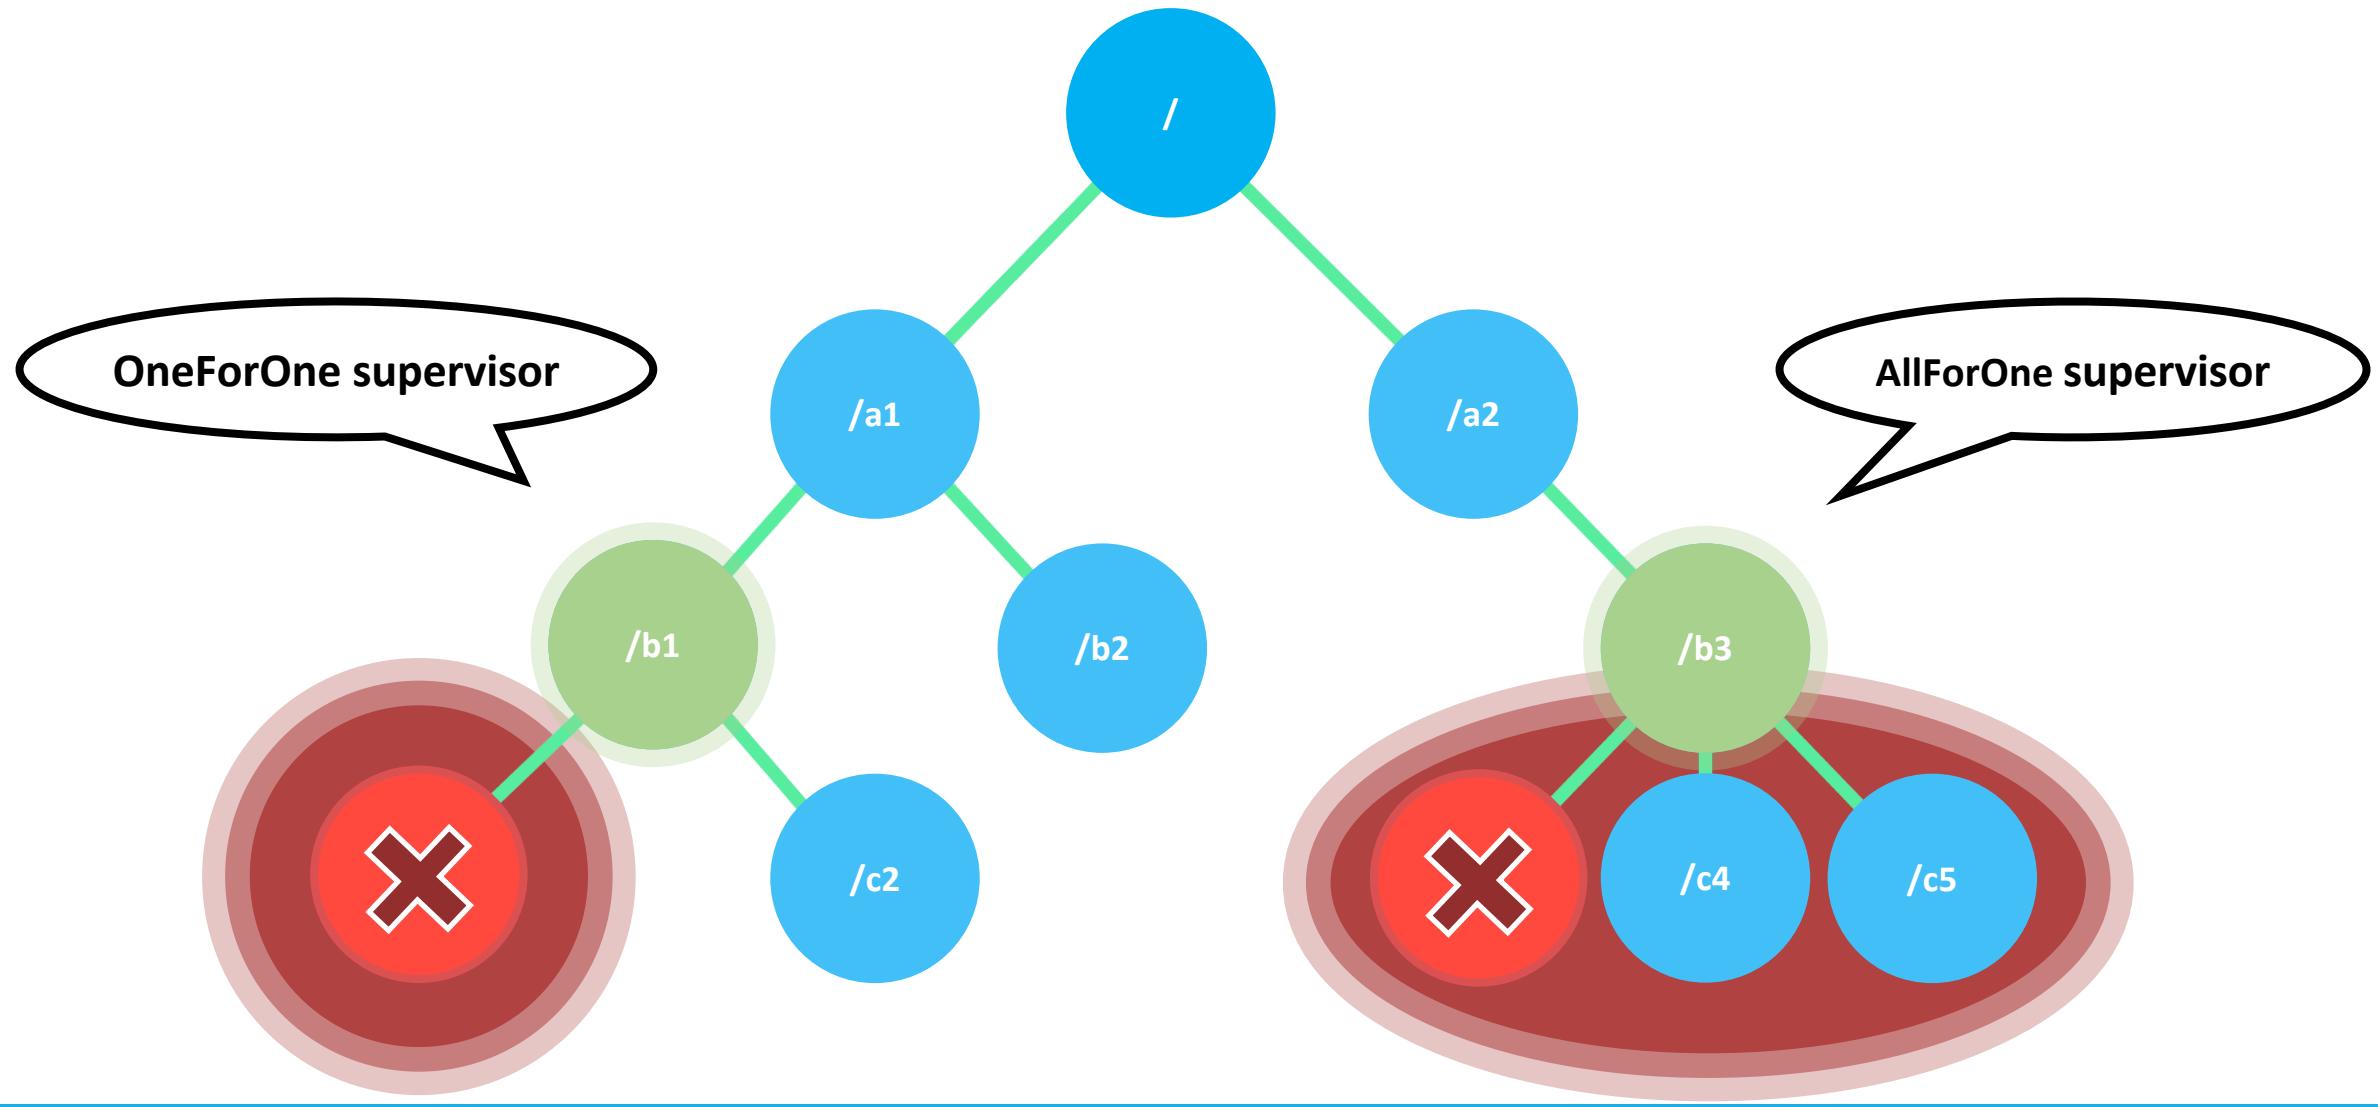


# Supervision



# Let it crash





# Actor Supervision

```
protected override SupervisorStrategy SupervisorStrategy()
{
    return new OneForOneStrategy(
        maxNrOfRetries: 10,
        withinTimeRange: TimeSpan.FromMinutes(1),
        localOnlyDecider: ex =>
    {
        if (ex is ArithmeticException)
        {
            return Directive.Resume;
        }

        return
Akka.Actor.SupervisorStrategy.DefaultStrategy.Decider.Decide
(ex);
    });
}
```

Akka.Game

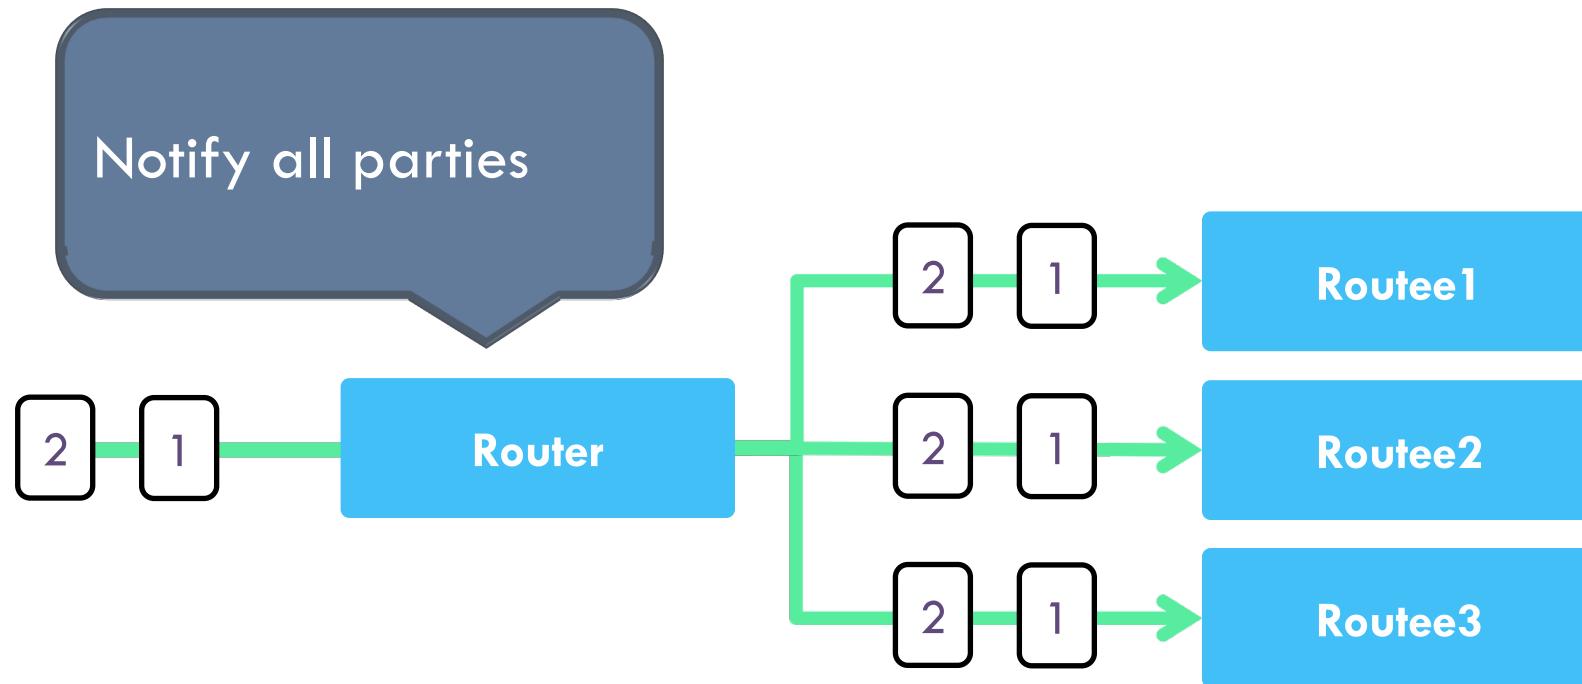


# Routing



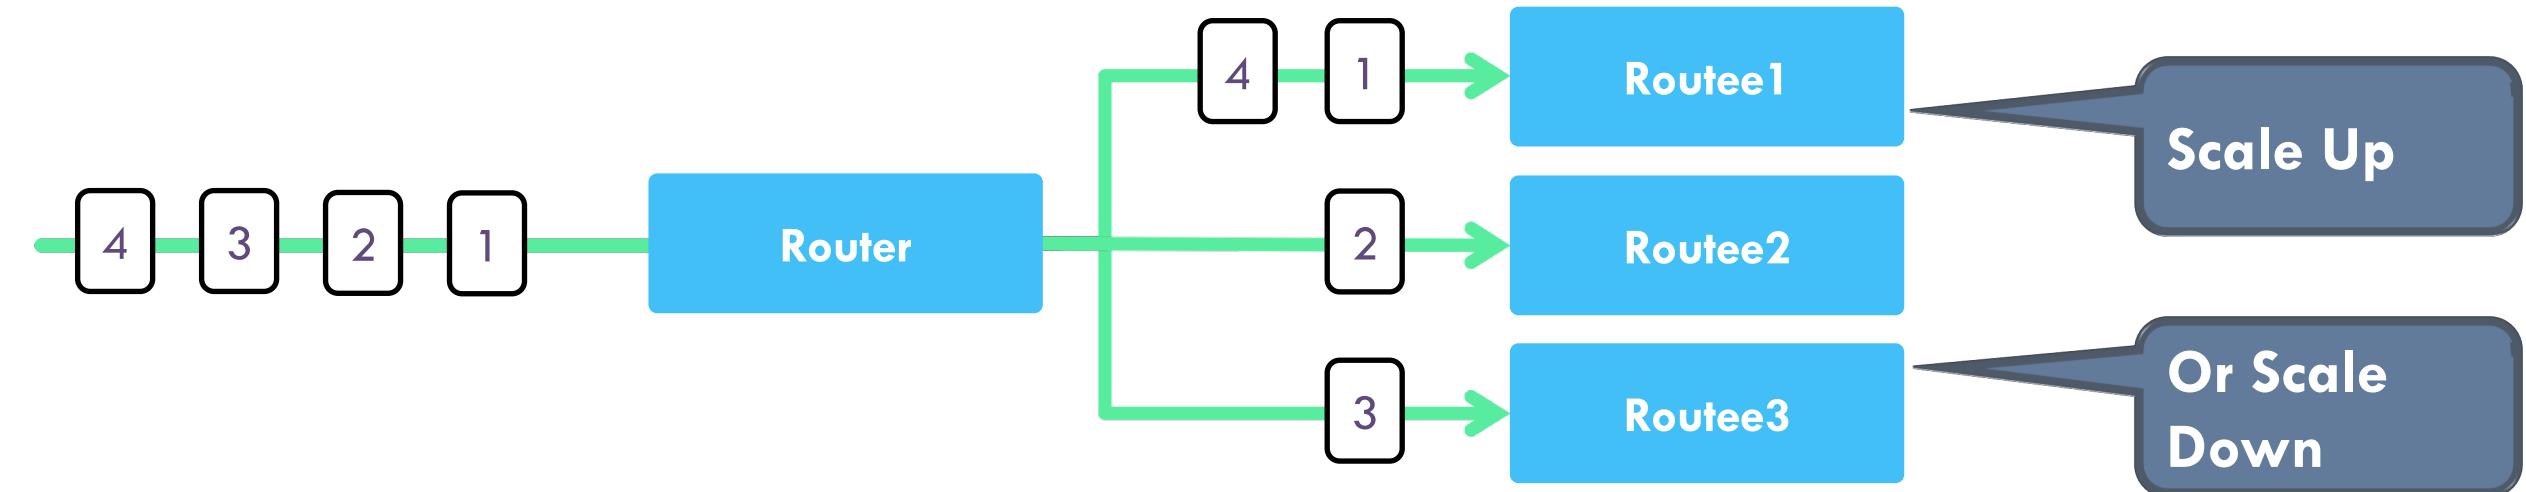
# Akka.Net Routing Strategies

[Broadcast](#) router will as the name implies, broadcast any message to all of its routees



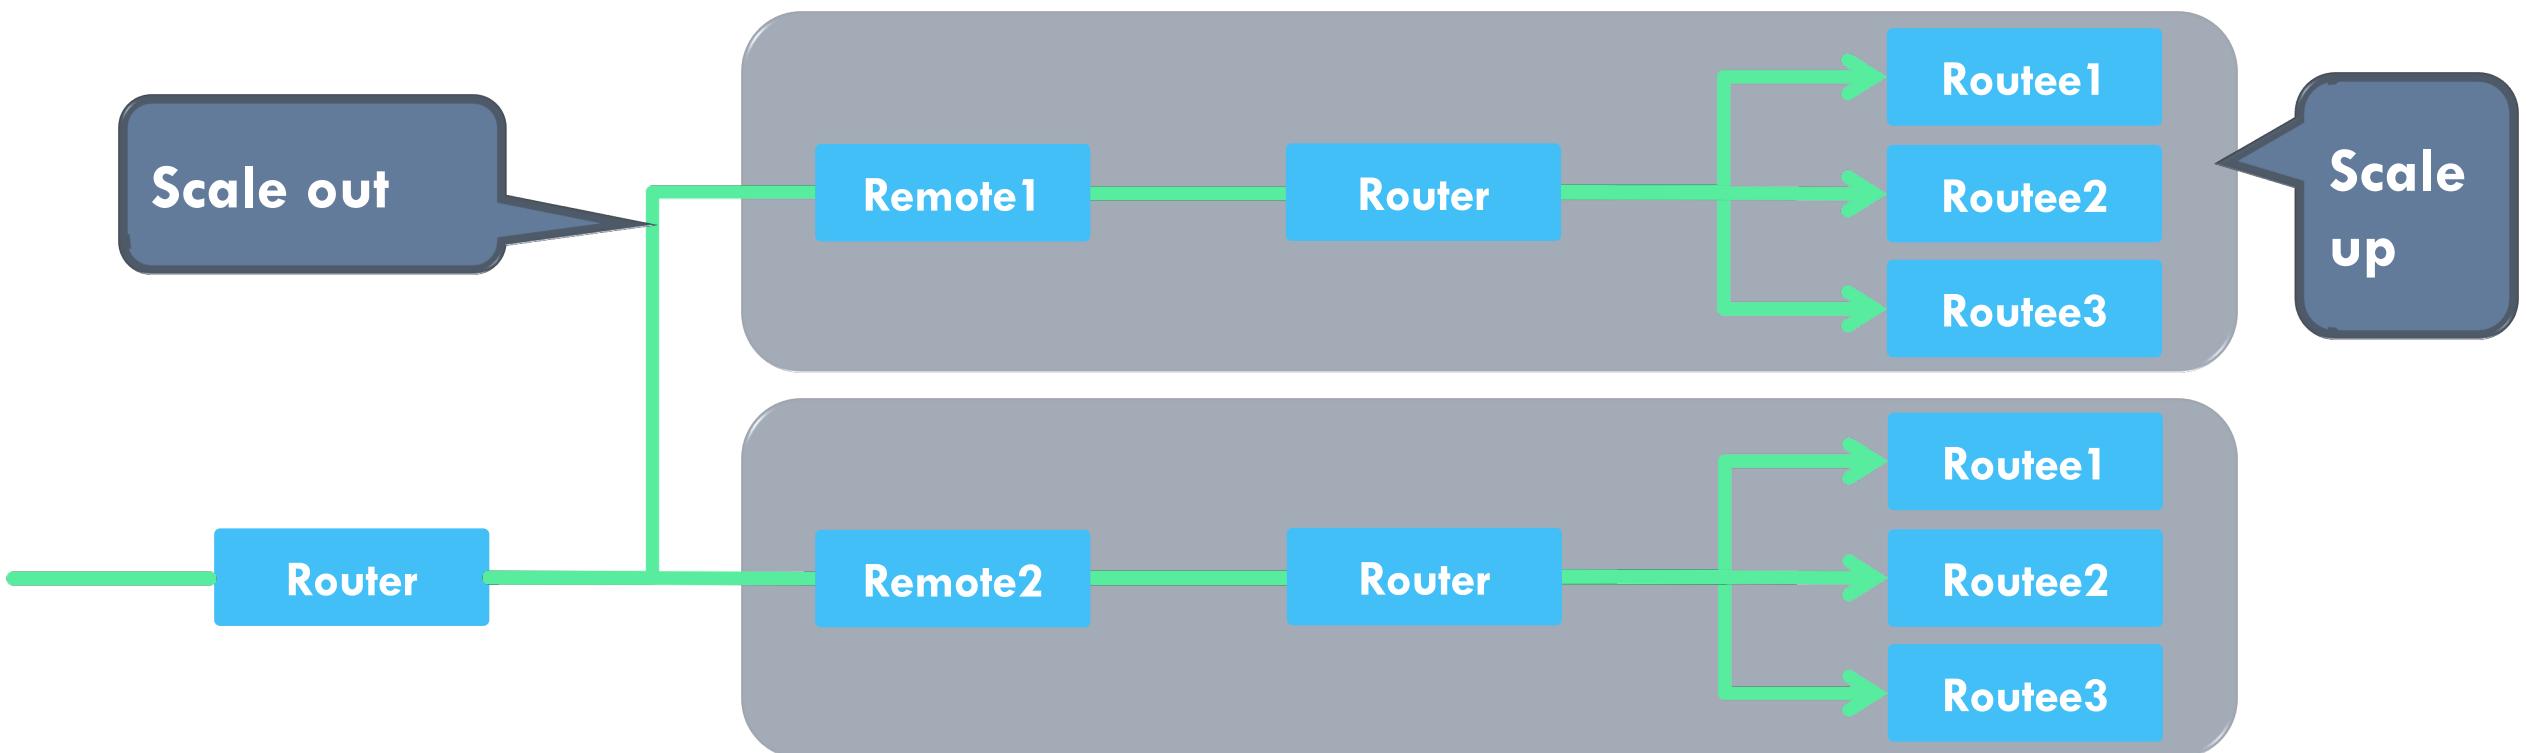
# Akka.Net Routing Strategies

[RoundRobin-Pool](#) router uses round-robin to select a connection. For concurrent calls, round robin is just a best effort.



# Akka.Net Routing Strategies

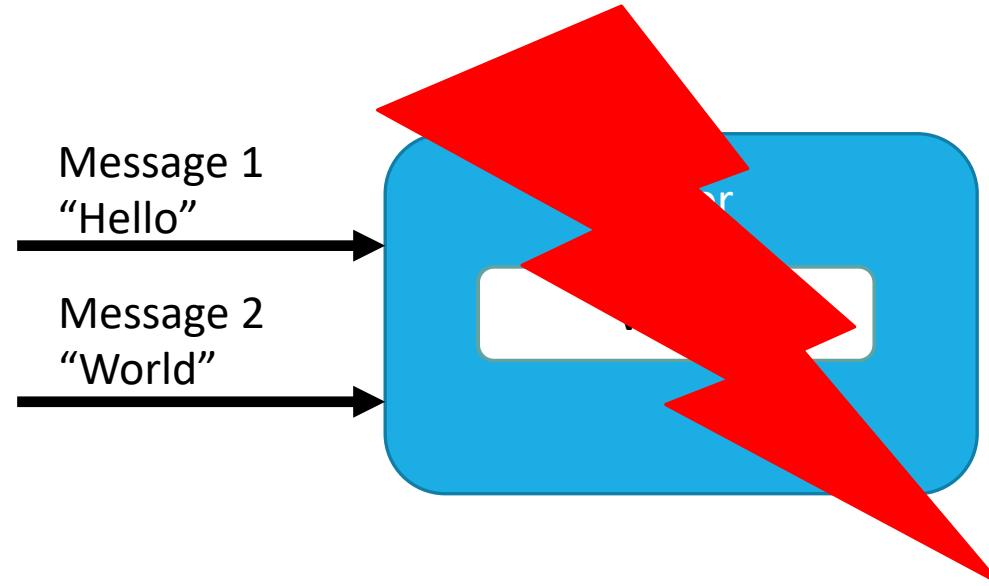
[RoundRobin-Group](#) router uses round-robin to select a connection. For concurrent calls, round robin is just a best effort.



# Actor Persistence

# Actor persistence

---

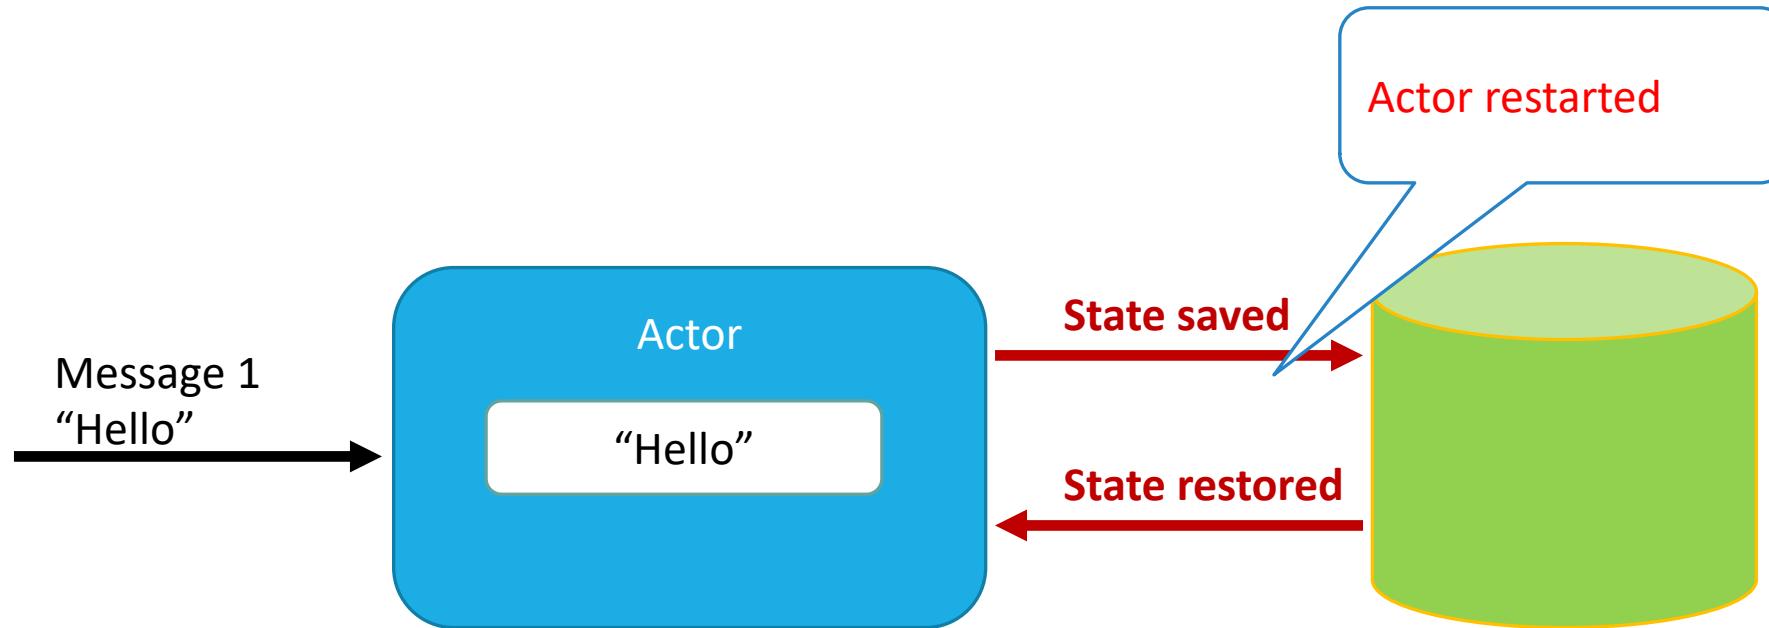


If for example there is an internal exception, the process crash and force the server and actor to restart

**The internal state is lost**

# Actor persistence

---



# Supported Persistence Stores

- Sql
- Postgres
- Mongo
- Redis



akka.persistence {  
 journal {  
 mongodb {  
 # qualified type name of the MongoDB persistence journal actor  
 class = "Akka.Persistence.MongoDb.Journal.MongoDbJournal, Akka.Persistence.MongoDb"  
  
 # connection string used for database access  
 connection-string = ""  
 # Should corresponding journal table's indexes be initialized automatically?  
 auto-initialize = off  
  
 # dispatcher used to drive journal actor  
 plugin-dispatcher = "akka.actor.default-dispatcher"  
  
 # MongoDB collection corresponding with persistent journal  
 collection = "EventJournal"  
  
 # metadata collection  
 metadata-collection = "Metadata"  
 }  
 }  
}

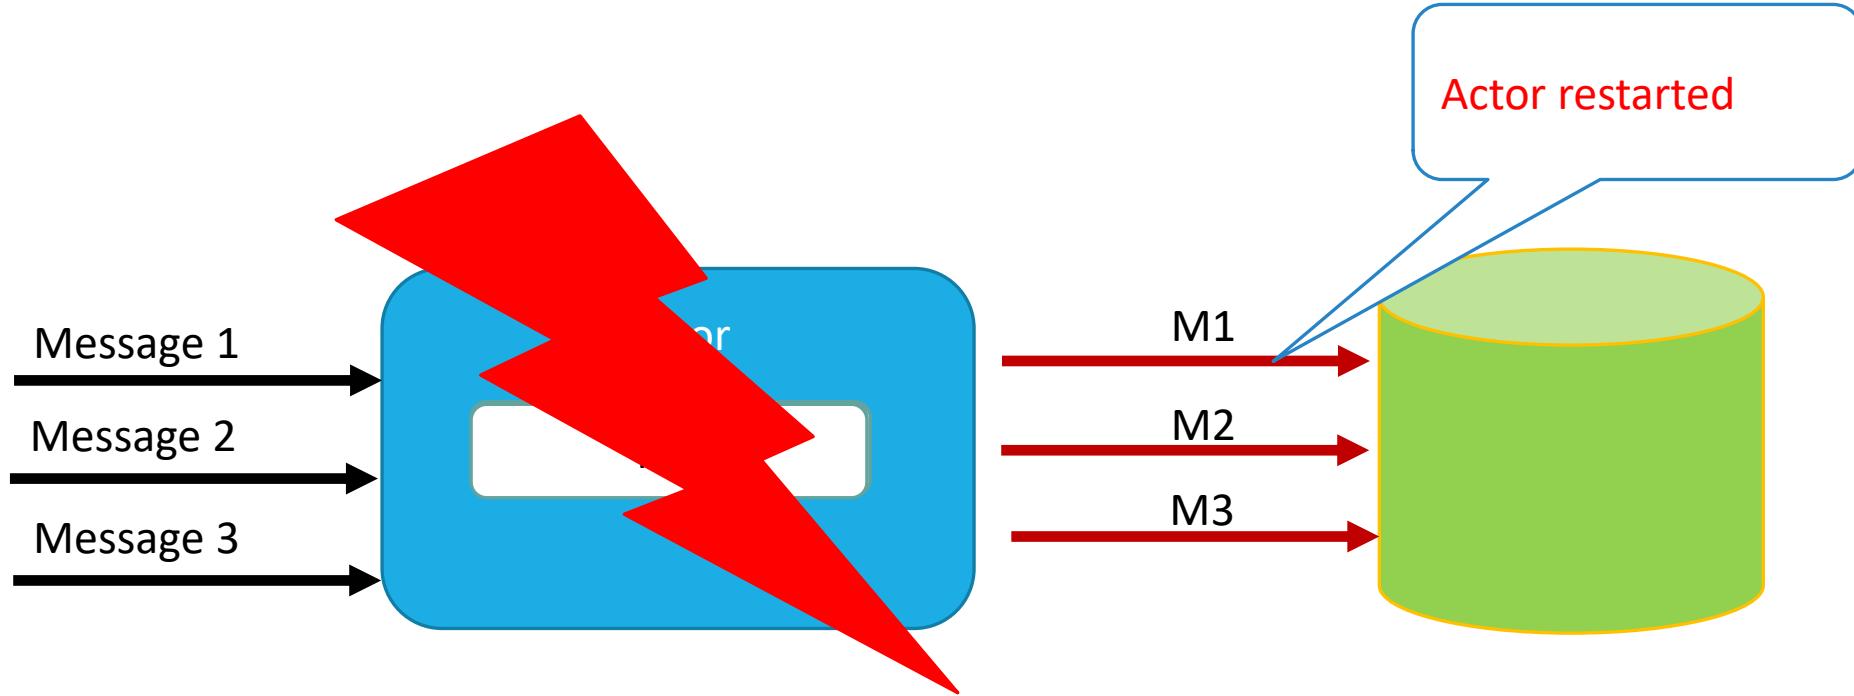
```
-> % docker run --name mongodb -p 27017:27017 -d mongo
```

snapshot-store {  
 mongodb {  
 # qualified type name of the MongoDB persistence snapshot actor  
 class = "Akka.Persistence.MongoDb.Snapshot.MongoDbSnapshotStore, Akka.Persistence."  
  
 # connection string used for database access  
 connection-string = ""  
 # Should corresponding snapshot's indexes be initialized automatically?  
 auto-initialize = off  
  
 # dispatcher used to drive snapshot storage actor  
 plugin-dispatcher = "akka.actor.default-dispatcher"  
  
 # MongoDB collection corresponding with persistent snapshot store  
 collection = "SnapshotStore"  
 }  
}

<https://github.com/AkkaNetContrib/Akka.Persistence.MongoDB>

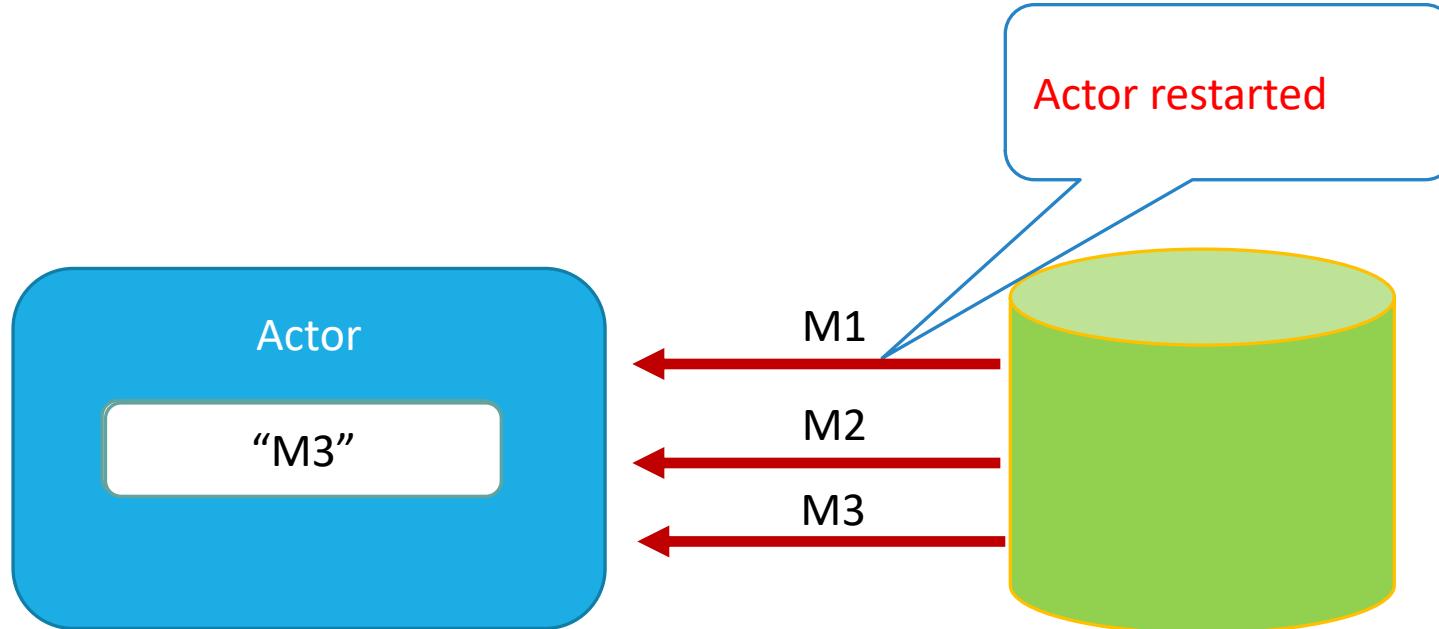
# Actor persistence

---



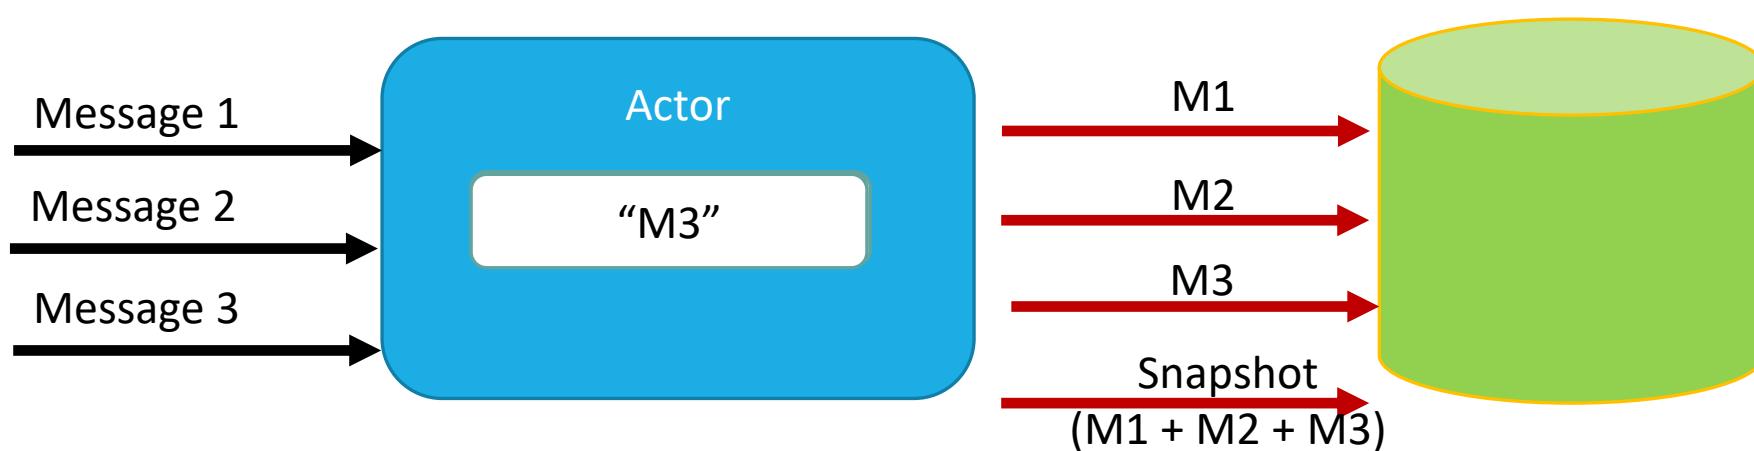
# Actor persistence

---

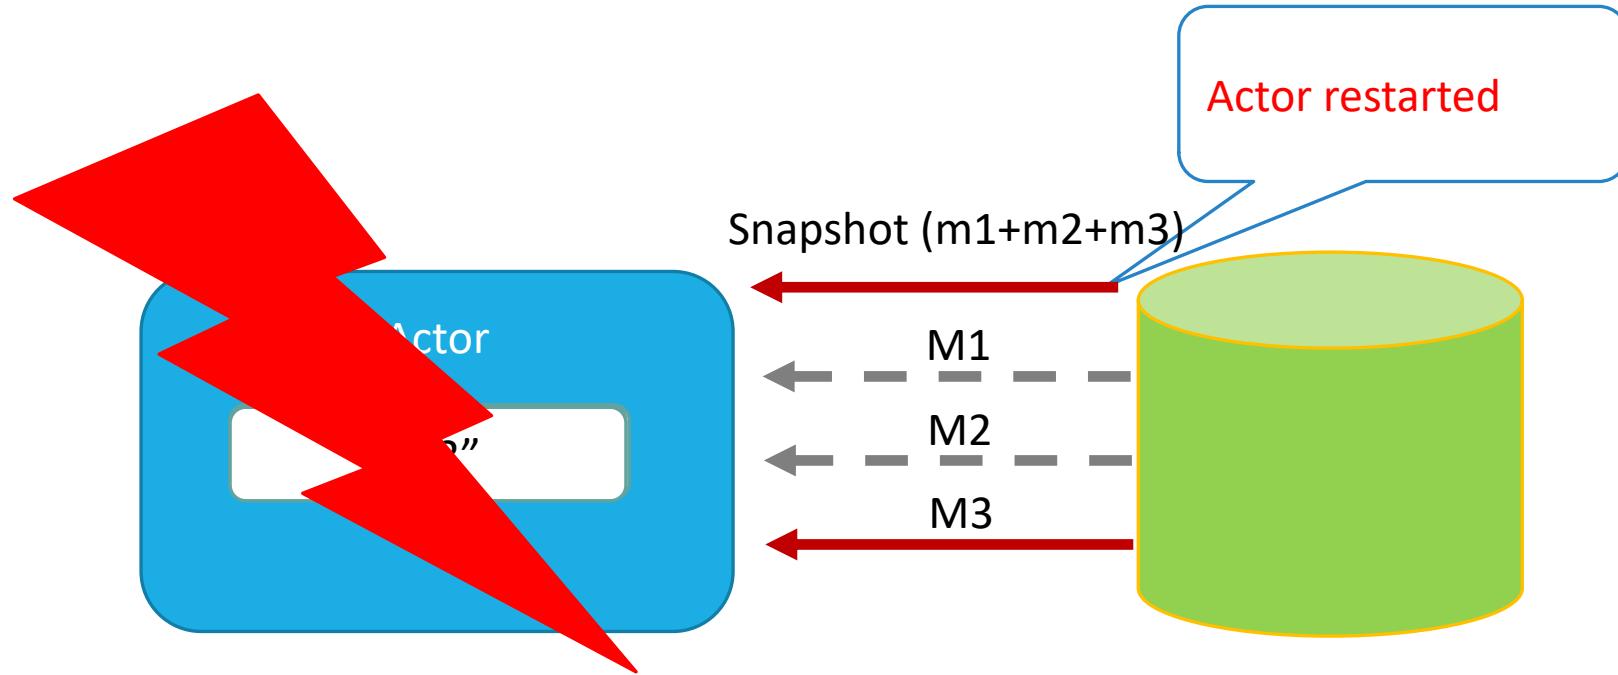


# Actor persistence - Snapshot

---



# Actor persistence - Snapshot



# Remoting



# *Location Transparency*



Protocol      Actor System      Address      Path

**Akka.tcp://System@10.55.1.25:9234/actorName**

# Remotely Deploying Actors

---

Deploying an actor means two things simultaneously:

Creating an actor instance with specific, explicitly configured properties

Getting an ActorRef to that actor

```
using (var system = ActorSystem.Create("Deployer", ConfigurationFactory.ParseString(@" // CLIENT
    akka { ... CONFIG AS BEFORE ... }
        remote {
            helios.tcp {
                port = 0
                hostname = localhost
            }
        }
    }"))
{
    var remoteEcho1 = system.ActorOf(Props.Create(() => new EchoActor()), "remoteecho");

    var echoActor = system.ActorOf(Props.Create(() => new HelloActor(remoteEcho1)));
    echoActor.Tell("Hello")
```

# Remote deployment configuration

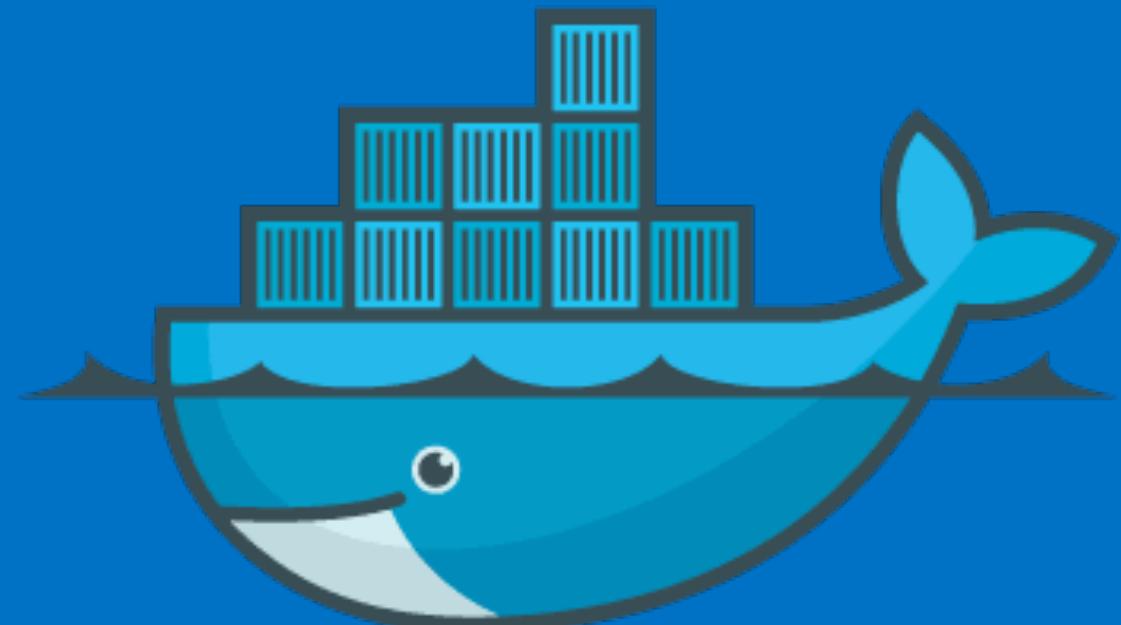
Human-Optimized Config Object Notation is a configuration format that aims to provide extensibility and flexibility, while at the same time making configuration values easy to edit by humans

## Hocon configuration

```
akka {  
    actor {  
        provider = "Akka.Remote.RemoteActorRefProvider, Akka.Remote"  
        deployment {  
            /remoteactor {  
                router = round-robin-pool  
                nr-of-instances = 1  
                remote = "akka.tcp://RemoteSystem@0.0.0.0:8090"  
            }  
        }  
    }  
    remote {  
        dot-netty.tcp {  
            port = 8091  
            hostname = 127.0.0.1  
        }  
    }  
}
```

## Remote node

```
let main args =  
use system = System.create "remote-system" config  
System.Console.ReadLine()  
0
```



docker

# What is Docker ?

---



- Lightweight, open, secure platform
- Simplify building, shipping, running apps
- Shipping container system for code
- Runs natively on Linux or Windows Server
- Runs on Windows or Mac Development machines (with a virtual machine)
- Relies on "images" and "containers"

# Dockerfile

---

```
FROM mono:latest  
  
MAINTAINER Riccardo Terrell  
  
ENV CODEDIR /var/code  
RUN mkdir -p $CODEDIR  
copy . $CODEDIR  
WORKDIR $CODEDIR  
  
RUN chmod +x $CODEDIR  
RUN xbuild ./ActorRemote.sln /property:Configuration=Release  
  
EXPOSE 9234  
  
ENTRYPOINT ["mono", "./AkkaExamples/bin/Release/AkkaExamples.exe"]
```



Dockerfile



Docker Image

# Docker-Compose

```
version: "3.7"
services:
  webapp:
    build:
      context: ./dir
      dockerfile: Dockerfile-alternate
      args:
        buildno: 1
    volumes:
      - type: volume
        source: mydata
        target: /data
        port: '8090:8095'
    volume:
      nocopy: true
```

# Akka.Fractal



# Clustering

- Load balancing
- Fault-tolerant
- Scalability

# Actor Clustering

---

*“Anything that can go wrong, will go wrong”*  
-- Murphy’s Law

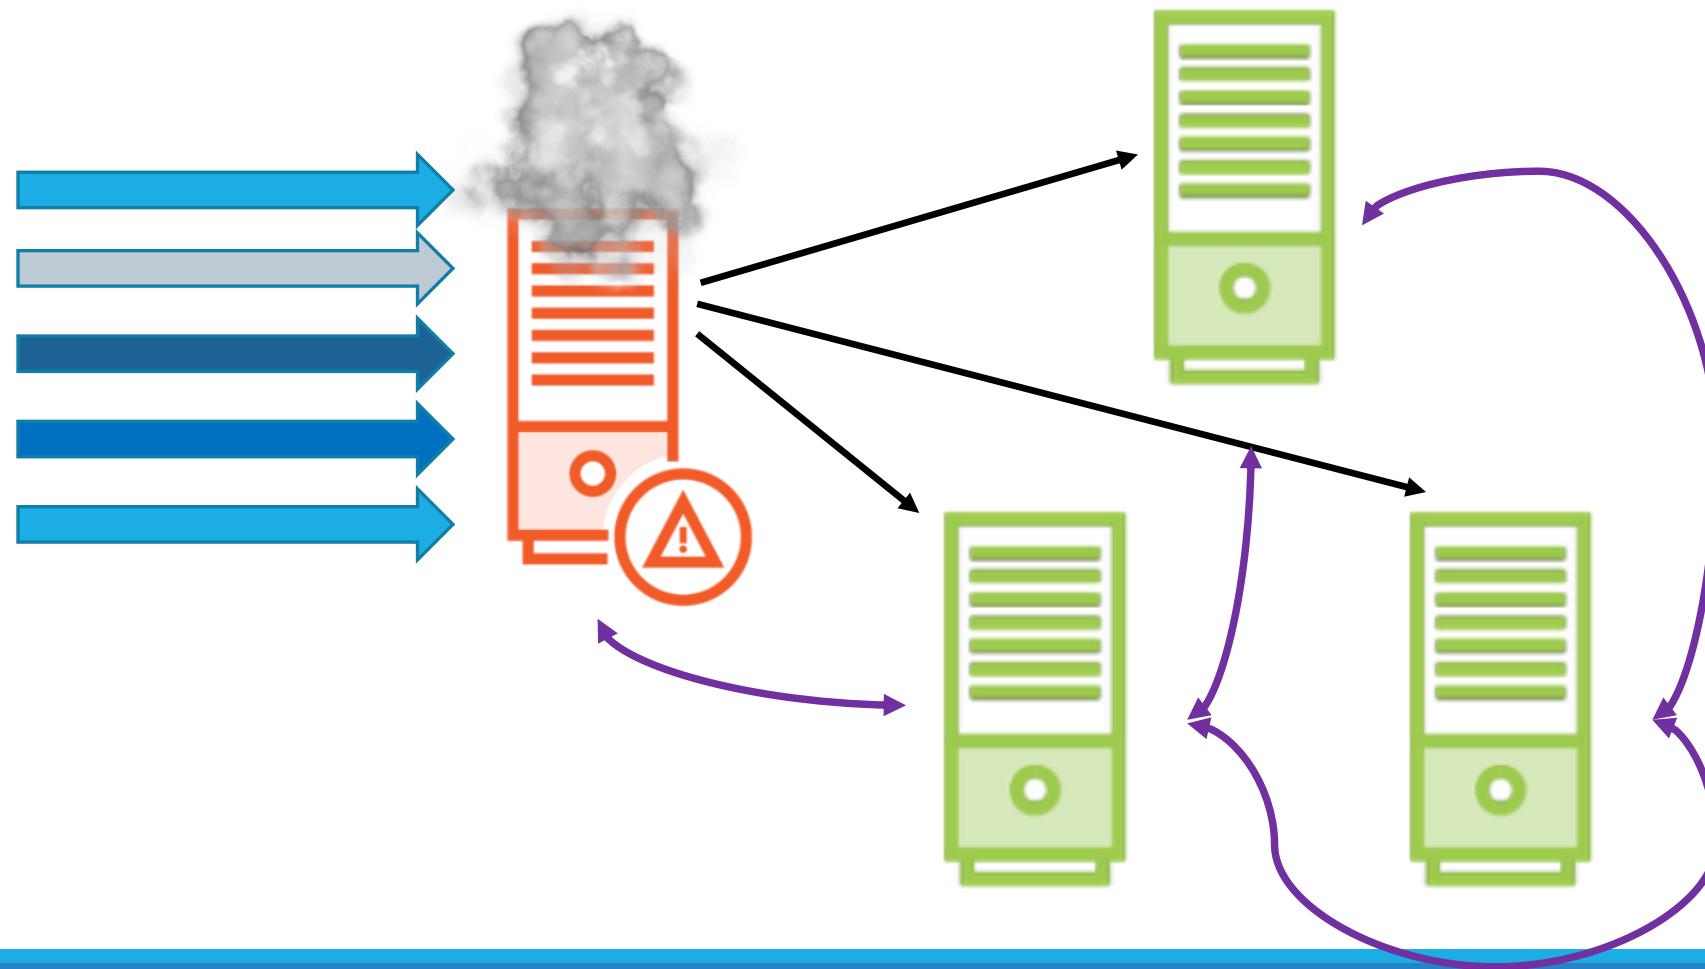


# Cluster Terminologies

- **Node**: a logical member of a cluster.
- **Cluster**: a set of nodes joined through the membership service. Multiple Akka.NET applications can be a part of a single cluster.
- **Gossip**: underlying messages powering the cluster itself.
- **Leader**: single node within the cluster who adds/removes nodes from the cluster.
- **Role**: a named responsibility or application within the cluster. A cluster can have multiple Akka.NET applications in it, each with its own role. A node may exist in 0+ roles simultaneously.
- **Convergence**: when a quorum (simple majority) of gossip messages agree on a change in state of a cluster member.

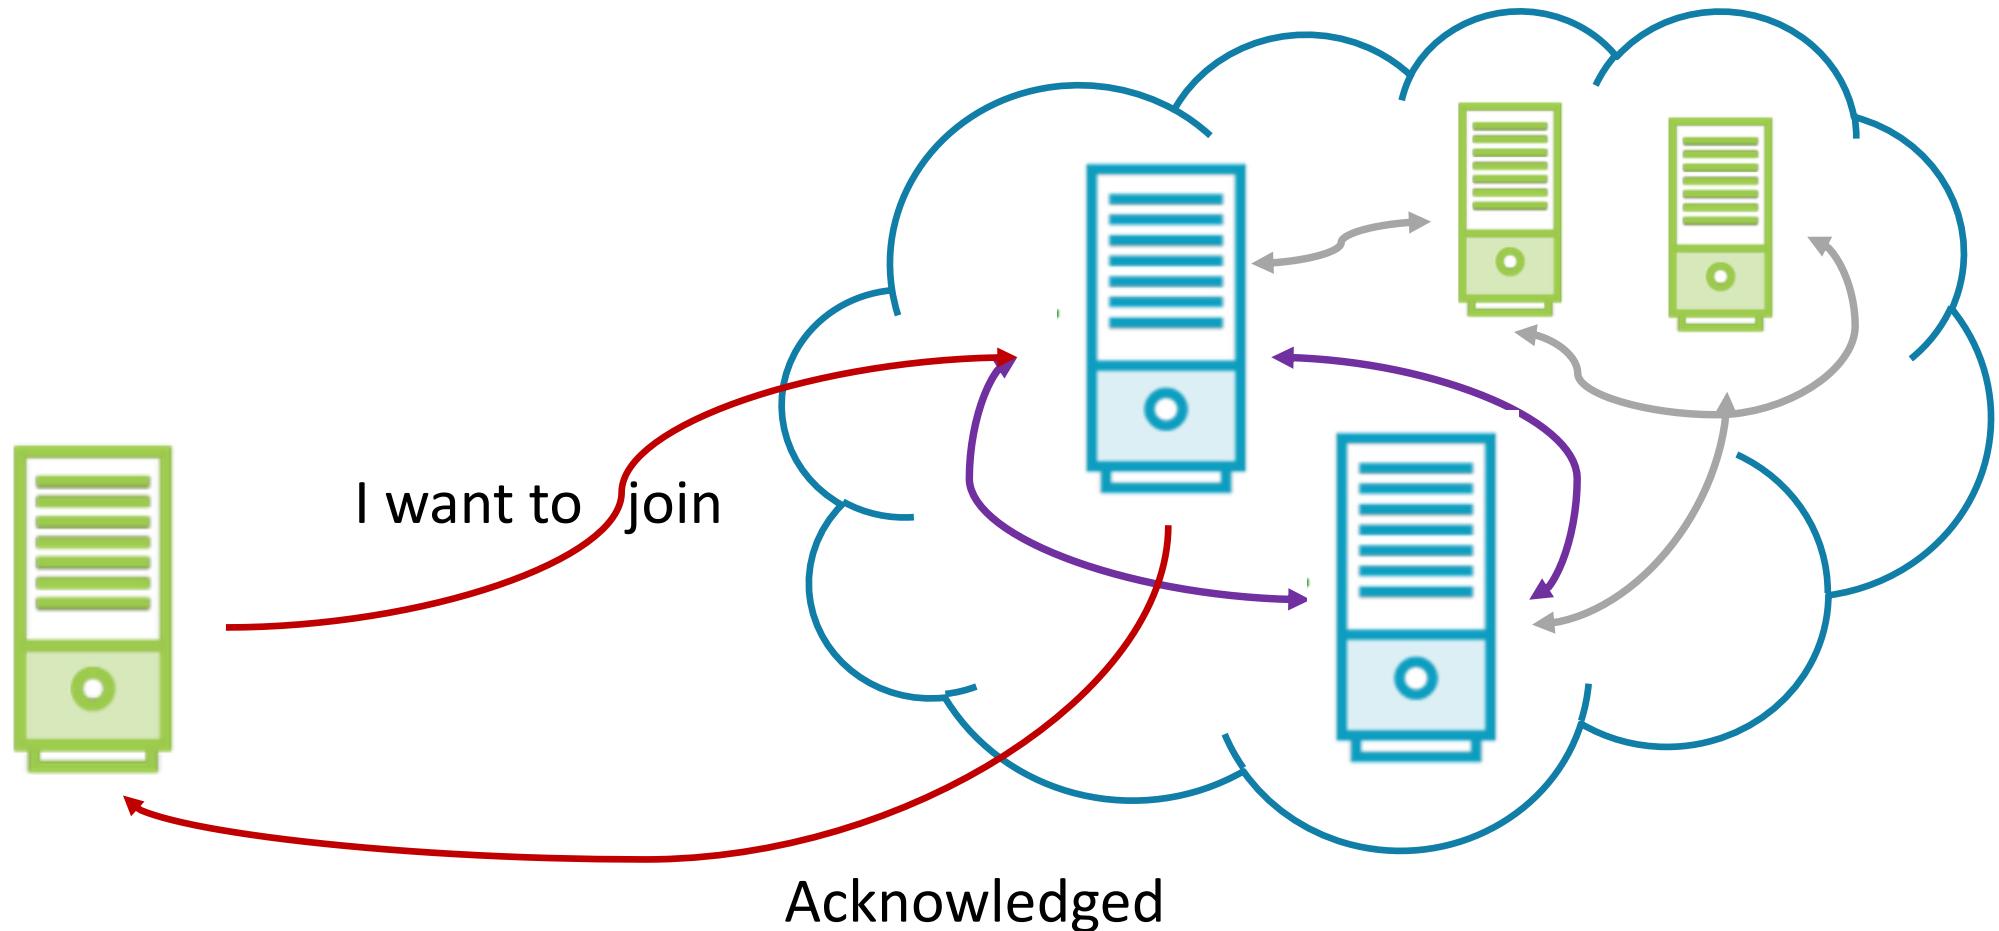
# Actor Clustering

---



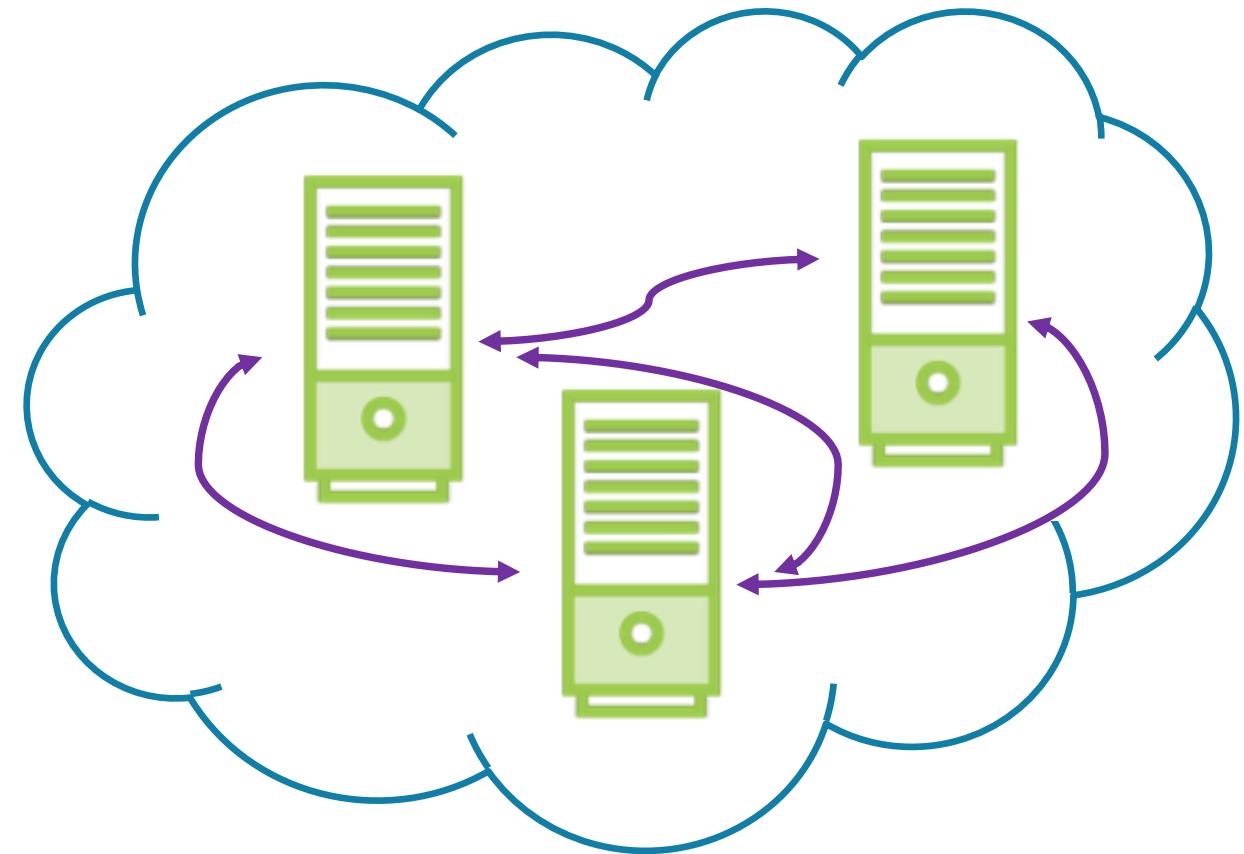
# Joining the Actor Cluster

---



# Joining the Actor Cluster

---

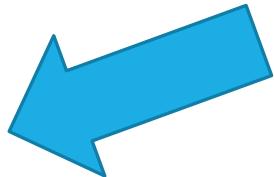


# Gossip



# Clustering – Just a configuration matter

```
let workerSystem = System.create "cluster" <| Configuration.parse """
akka {
    actor {
        provider = "Akka.Cluster.ClusterActorRefProvider, Akka.Cluster"
    }
}
remote {
    helios.tcp {
        hostname = "localhost"
        port = 0
    }
}
cluster {
    seed-nodes = [ "akka.tcp://cluster@localhost:8091/" ]
}
"""
```



# Clustering – Just a configuration matter

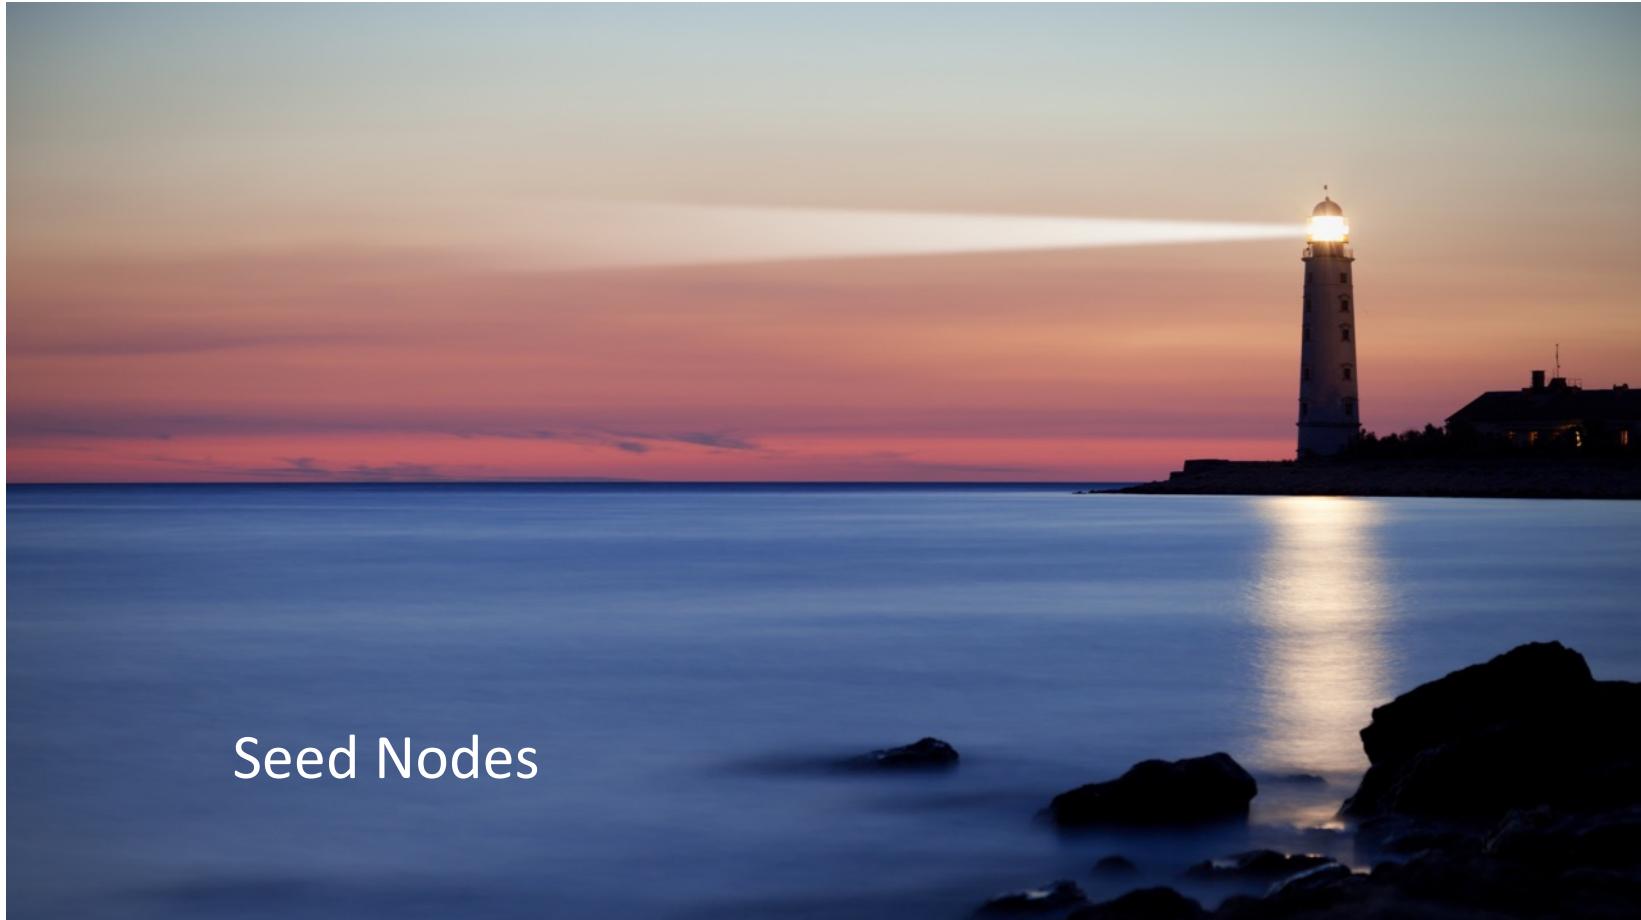
```
let masterSystem = System.create "cluster" <| Configuration.parse """  
akka {  
    actor {  
        provider = "Akka.Cluster.ClusterActorRefProvider, Akka.Cluster"  
    }  
}  
remote {  
    helios.tcp {  
        hostname = "localhost"  
        port = 8091  
    }  
}  
cluster {  
    seed-nodes = [ "akka.tcp://cluster@localhost:8091/" ]  
}  
}
```



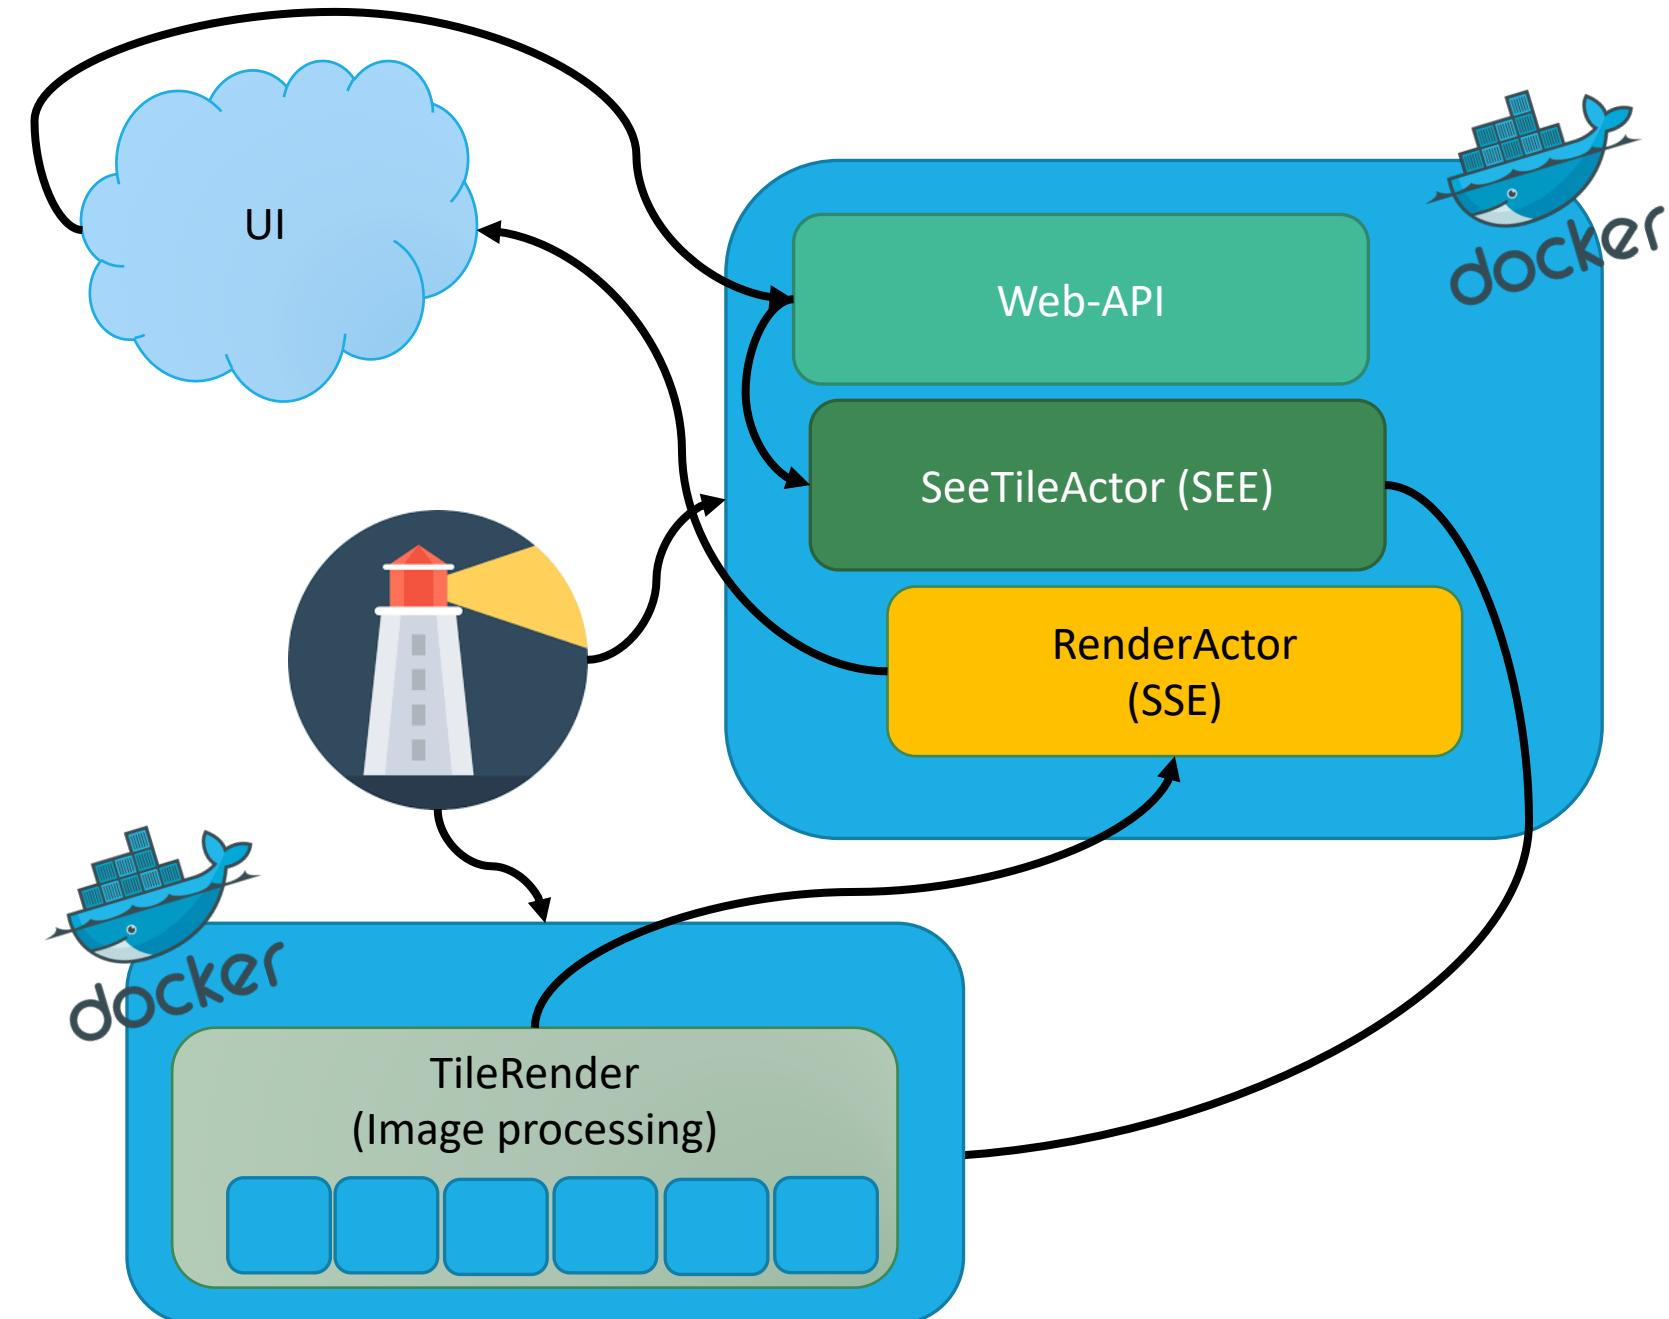
# Lighthouse

Lighthouse dedicated is a simple but effective seed node

Seed Nodes



# Project Remote Distributed Image (fractal) Processing



# Summary

---

- Message passing programming model allows to easily scale up and/or scale out your application
- Supervision allows to build fault tolerant and resilient systems
- You can distribute the work and deploy actor cross network using remoting
- Routing provides different work distribution strategies to maximize/optimize the resources consumption
- Clustering allows the collaboration of different Actor-Systems to work as single unit, enabling a system to be responsive, resilient, elastic and message driven.

# contacts

*Source*

<https://github.com/rikace/raleighcc2019>

*Twitter*

@trikace

*Blog*

[www.rickyterrell.com](http://www.rickyterrell.com)

*Email*

[tericcardo@gmail.com](mailto:tericcardo@gmail.com)

*Github*

[www.github.com/rikace/](https://www.github.com/rikace/)



*The tools we use have a profound (and devious!) influence on our thinking habits, and, therefore, on our thinking abilities.*

-- Edsger Dijkstra

That's all Folks!